

## Assignment 3: Programming with MPI

### (Due Wednesday, 5/18/16)

This assignment is to practice message-passing programming with the use of the MPI library. You'll implement an external sorting algorithm, *i.e.*, reading data from a file, sorting the data, and write the result to another file.

For this assignment, all students will implement the same program. However, CS515 students are required to include timing results for the program (see below). This assignment carries a total of 10 points.

Download and unzip the file `assign3.zip`, you'll see a directory `assign3` containing some program files.

### External Sorting

Your program should be called `extsort.c`. It should be implemented in the SPMD style. It should take two command-line arguments: the input file name and the output file name. The program should read data from the input file; follow the sorting algorithm shown below to sort the data; and write the result to the output file.

### Algorithm

The sorting algorithm is a simplified version of sample sort on integers. Assume data size is  $N$  and the total number of processes is  $P$  ( $P \geq 2$ ). (Assume also that  $N > 10P$ .)

1. Process 0 reads in all data from the input file.
2. Process 0 sorts the first  $10P$  elements, and selects elements at positions 10, 20, ..., and  $10(P - 1)$  as pivots. (They will be referred to as `pivot[0]`, `pivot[1]`, ..., and `pivot[P - 1]`.)
3. Process 0 partitions the data into  $P$  buckets — elements whose values are smaller than `pivot[0]` are placed in `bucket[0]`, elements whose values are in between of `pivot[0]` and `pivot[1]` are placed in `bucket[1]`, and so forth.
4. Process 0 keeps `bucket[0]` to itself, and sends the rest  $P - 1$  buckets to their corresponding processes, *i.e.* `bucket[i]` to process  $i$ .
5. Every process sorts its bucket using quicksort.
6. The processes write their results to the output file, in the process rank order.

A copy of the quicksort program can be found in the `assign3` directory. You may copy the useful part into your `extsort.c` program.

### Data File Format

The data to be sorted are four-byte integers (C type `int`). Both the input and output files are *byte* files in which each consecutive group of four bytes encode an integer. For example, for the four integers, 860, 386, 103, and 282, the data file will contain the following 16 bytes:

```
5c 03 00 00 // 1st int 860
82 01 00 00 // 2nd int 386
67 00 00 00 // 3rd int 103
1a 01 00 00 // 4th int 282
```

Note that the content of a binary file is not directly viewable. To see a binary content, use the Linux utility, `od`, with a proper switch:

```
linux> od -i data1k -- display binary content as integers
```

A pair of programs, `datagen.c` and `verify.c`, are provided to you for dealing with the data files. The program `datagen` takes an integer command-line argument, `N`, and generates a random permutation of `N` integers, `1, ..., N`, which can be saved in a data file:

```
linux> ./datagen 1024 > data1k
```

The program `verify` can be used to verify that the integer values in a given data file are sorted in an ascending order:

```
linux> ./verify out1k
Data in out1k are sorted.
```

## File I/O

Use MPI's file I/O routines to handle input and output. For the program's input, only Process 0 is involved. So you may use `MPI_COMM_SELF` when opening the input file, to restrict access to itself. Also there is no need to call the routine `MPI_File_set_view()`.

For the output, all processes are involved. Here you have two choices.

- (1) You can arrange the processes to take turn to access the output file. Each process opens and closes the file for its own use. (Hence use `MPI_COMM_SELF` again.) It appends its output to the end of the file. For this approach, the issue to resolve is to have processes take turns in the process rank order. (*Hint*: Think about the `ring.c` program, and use messages to enforce the required order.)
- (2) Have all processes write to the same file concurrently. For this approach to work, you need to have all processes call `MPI_File_open()` with `MPI_COMM_WORLD`, and use the routine `MPI_File_set_view()` to set their offset values. Reference the `iodemo.c` program for the usage pattern of this routine. One issue you need to resolve is to figure out the right offset for each process.

## [CS515 Students] Timing

Insert the MPI timing routine `MPI_Wtime()` in your program to collect timing data. You may want to measure two versions of total elapsed time: one includes everything and the other excludes I/O actions. Find the right points to insert the timing routine calls for these measurements.

## Report

As usual, write a summary report on the assignment (around 2 pages). You may include anything you want to discuss. If you have collected timing results, include them with some comments.

## What to Turn In:

Make a `zip` file containing your source program and the summary. Use the `Dropbox` on the D2L site to submit.