

Introduction to Chapel

Jingke Li

Portland State University

Chapel

A new high-level parallel programming language. Started out as a Cray project with DARPA funding, now is open-source.

- ▶ Support multiple parallel programming paradigms
- ▶ Modern, comprehensive, and programmer friendly
- ▶ Not fully developed, performance is not yet competitive

Topics (for now):

- ▶ Basic language features
- ▶ Data parallel features
- ▶ Task parallel features

Disclaimer: Some materials are adapted from Cray's tutorials on Chapel.

The “Hello, World!” Program

- ▶ Ver 1. Quick, scripting-language style.

```
writeln("Hello, world!");
```

- ▶ Ver 2. With module declaration.

```
module Hello {  
  writeln("Hello, world!");  
}
```

Module-level code is executed before program starts.

Note: These two versions are exactly equivalent. Chapel compiler uses an inference engine to fill in the missing parts in Ver 1 to convert it to Ver 2.

The “Hello, World!” Program (cont.)

- ▶ Ver 3. Full, structured programming style.

```
module Hello {  
  proc main() {  
    writeln("Hello, world!");  
  }  
}
```

This version is slightly different from the previous versions — the “main” procedure is executed when program begins running.

Data Types

- ▶ Primitive Types:
bool, int, uint, real, image, complex, string
- ▶ Type's bit width can be explicitly specified:
int(16), int(32), real(32)
- ▶ Type casts are allowed, but with a different syntax from C:

```
5:int(8)           // store value as int(8) rather than int  
"54":int           // convert string to an int  
249:complex(64)    // convert int to complex(64)
```
- ▶ Every primitive type has a default value.

Variables and Constants

- ▶ Compile-time constant: param
- ▶ Execution-time constant: const
- ▶ Execution-time variable: var

```
param pi: real = 3.14159;  
const debug = true;           // inferred to be bool  
var count: int;               // initialized to 0
```

Chapel compiler provides initial value if user does not provide one. It also infers type if the type info is missing.

Note: Any of these items can be declared with config to make them configurable.

The “Hello, World!” Program (again)

- Ver 1a. With a configurable message.

```
config const message = "Hello, world!";
writeln(message);
```

```
linux> ./hello2
Hello, world!
linux> ./hello2 --message="Hi!"
Hi!
```

However, you'll get a surprising result if you type

```
linux> ./hello2 --message="Hi!!!"
```

— because "!!!" is interpreted as a shell-command.

Tuples and Arrays

- Tuples provide lightweight grouping of values:

```
var coord: (int, int, int) = (1, 2, 3);
var coord2: 3*int = coord;
var (i, j, k) = coord;
var triple: (int, string, real) = (7, "eight", 9.0);
```

- Arrays have flexible syntax forms:

```
var A[1..3] int = [5,6,7];
var B: [1..3, 1..5] real;
var C: [1..3][1..5] real;
var i = 1, j = 2;
var ij = (i,j);
B[ij] = 1.0;
B[ij] = 2.0;
B[i,j] = 3.0;
B(i,j) = 4.0;
```

For Loops

- Flexible Syntax — with or without the keyword do; different forms of index domain:

```
var A: [1..3] string = ["April", "May", "June"];
for i in 1..3 do
  write(A(i));
for a in A {
  a += ", 2014";
}
write(A);
```

- Compressed nested-loop form:

```
var B: [0..9] real;
for (b,i,j) in zip(B, 1..10, 2..20 by 2) do
  b = j + i/10.0;
writeln(B);
// 2.1 4.2 6.3 8.4 10.5 12.6 14.7 16.8 18.9 21.0
```

Iterators

- Generalized form of loops; requires multi-threading support.

```
iter fibonacci(n) {
  var current = 0, next = 1;
  for 1..n {
    yield current; // return current fib number
    current += next; // create a new fib number
    current <=> next; // swap the two values
  }
}

for f in fibonacci(7)
  do writeln(f);
```

Procedures

- Can set parameter's default value

```
proc writeCoord(x: real = 0.0, y: real = 0.0) {
  writeln((x,y));
}
```

- Can supply arguments out of order

```
writeCoord(2.0); // (2.0, 0.0)
writeCoord(y=2.0); // (0.0, 2.0)
writeCoord(y=2.0, 3.0); // (3.0, 2.0)
```

- I/O procedures have flexible parameter forms

```
write(expr-list) // prints an expr list
read(expr-list) // reads into a param list
read(type-list) // returns a tuple
```

Records and Classes

Both may contain fields and variables.

The difference: records are value-based, classes are reference-based.

- Methods without arguments need not use parenthesis

```
class circle { var radius: real; }
proc circle.circumference {
  return 2 * pi * radius;
}
var c1 = new circle(radius=1.0);
writeln(c1.circumference);
```

- Methods can be defined for any type

```
proc int.square() {
  return this**2;
}
writeln(5.square());
```

Chapel's Data-Parallel Features

Similar to Fortran 90/95, but more user friendly.

► Array Operations:

```
var a,b,c: [{1..5}] int; // declaring three int array vars
a = [i in {1..5}] i;    // a = [1,2,3,4,5]
b = 1;                 // b = [1,1,1,1,1]
c = a + b;              // c = [2,3,4,5,6]
writeln(c);
```

► Forall Loops:

```
var a: [{1..5}] int = 2;
forall i in {1..5} do
  a[i] += 1;
writeln(a);
```

Loop body instances are intended to be executed in parallel; but they must also be executable sequentially.

Domains

In Chapel, index sets can be manipulated separately from arrays — they are called *domains*. They can be assigned to variables, passed around, and so forth. (In other words, domains are “first-class” objects.)

```
var D = {1..5}; // Declare a domain
var a,b,c: [D] int; // declaring three int array vars
a = [i in D] i;    // a = [1,2,3,4,5]
b = 1;            // b = [1,1,1,1,1]
c = a + b;         // c = [2,3,4,5,6]
writeln(c);
```

```
D = {2..4}; // Now, redefine domain D
writeln(c); // What happens?
```

Domains (cont.)

Domain serves a fundamental role for data parallelism in Chapel.

- Arrays are defined in terms of domains. (previous examples)
- Forall are also defined over domains:

```
var D = {1..5};
var a: [D] int = 2;
forall i in D do
  a[i] += 1;
writeln(a);
```

- Domain serves as basis for aligning multiple arrays.
- Domain can be distributed over locales to easily transform a data-parallel program into a distributed program.

“Hello, world!” (Revisit)

► Data Parallel Version:

```
config const numMessages = 100;
forall msg in 1..numMessages do
  writeln("Hello, world! (from iteration ",
    msg, " of ", numMessages, ")");
```

► Distributed Data Parallel Version:

```
use CyclicDist;
config const numMessages = 100;
const MessageSpace = {1..numMessages}
dmapped Cyclic(startIdx=1);
forall msg in MessageSpace do
  writeln("Hello, world! (from iteration ",
    msg, " of ", numMessages, " owned by locale ",
    here.id, " of ", numLocales, ")");
```

Reduce and Scan

► Reduce:

```
total = + reduce A;
bigDiff = max reduce [i in Inner] abs(A[i]-B[i]);
(minVal, minLoc) = minloc reduce zip(A, D);
```

► Scan:

```
var A, B, C: [1..5] int;
A = 1; // A: 1 1 1 1 1
B = + scan A; // B: 1 2 3 4 5
B[3] = -B[3]; // B: 1 2 -3 4 5
C = min scan B; // C: 1 1 -3 -3 -3
```

Chapel's Task-Parallel Features

Three Constructs:

- begin — for creating a dynamic task with an unstructured lifetime
- cobegin — for creating a related set of heterogeneous tasks
- coforall — for creating a fixed or dynamic # of homogeneous tasks

For cobegin and coforall, the created tasks will joint back to the parent.

"Hello, world!" (Again)

► Task Parallel Version 1:

```
begin writeln("Hello, world!");  
writeln("Good bye!");
```

► Task Parallel Version 2:

```
cobegin {  
  writeln("Hello, world! (1)");  
  writeln("Hello, world! (2)");  
}
```

► Task Parallel Version 3:

```
config const numTasks = here.numCores;  
coforall tid in 0..  
  numTasks do  
  writeln("Hello, world! (from task " + tid  
    + " of " + numTasks + ")");
```

Task Synchronization

Four Mechanisms:

- Sync variables — equivalent to locks
- Single variables — similar to signals
- Atomic variables — providing atomic accesses
- Sync statements — similar to waits

Sync Variables

- Stores full/empty state along with normal value (default to full at initialization)
- A read (by default) blocks until full, leaves empty
- A write (by default) blocks until empty, leaves full

```
var lock$: sync bool;  
  
lock$ = true;  
critical();  
var lockval = lock$;
```

Note: As a suggested convention, synchronization variables are named with a \$ at the end (although it's not required).

Example: Consumer/Producer with Bounded Buffer

```
var buff$: [0..  
  #buffersize] sync real;  
  
begin producer();  
consumer();  
  
proc producer() {  
  var i = 0;  
  for ... {  
    i = (i+1) % buffersize;  
    buff$[i] = ...;  
  }  
}  
  
proc consumer() {  
  var i = 0;  
  while ... {  
    i = (i+1) % buffersize;  
    ...buff$[i]...;  
  }  
}
```

Single Variables

- Similar to sync variable, but stays full once written

```
var signal$: single real;  
  
begin signal$ = master();  
begin worker(signal$);  
begin worker(signal$);
```

Note: The default behavior of synchronization and single variables can be changed by user.

Atomic Variables

- Supports operations on variable atomically (with respect to other tasks)

```
var count: atomic int, done: atomic bool;  
  
proc barrier(numTasks) {  
  const myCount = count.fetchAdd(1);  
  if (myCount < numTasks) then  
    done.waitFor(true);  
  else  
    done.testAndSet();  
}
```

Atomic Methods:

```
read():t          compareExchange(old:tnew:t):bool  
write(v:t)        waitFor(v:t)  
exchange(v:t):t   add(v:t)  
testAndSet()      fetchAdd(v:t)  
clear()
```

Sync Statements

- ▶ Waits for all dynamically-scoped begins to complete

```
sync {  
  for i in 1..numConsumers {  
    begin consumer(i);  
  }  
  producer();  
}
```

Chapel's Other Features

(To be covered at a future time.)

Support for message-passing programming:

- ▶ Domains
- ▶ Domain maps
- ▶ Locales