PRAM and Other Models

Jingke Li

Portland State University

Jingke Li (Portland State University) CS 415/515 PRAM and Other Models

The RAM Model

RAM = Random Access Machine

A commonly used model for sequential computation.

- ► Machine Model:
 - ▶ A processor operating under the control of a sequential program.
 - ▶ A memory with M cells (M can be unbounded).
- Basic Operations:
 - ▶ Read the processor reads a datum from an arbitrary location in memory into one of its internal registers.
 - ► Compute the processor performs an (arithmetic) operation on data in register(s).
 - ▶ Write the processor writes the content of one register into an arbitrary memory cell.
- Cost Model:

Each basic operation takes one time unit to execute.

Jingke Li (Portland State University) CS 415/515 PRAM and Other Models

RAM Algorithm Examples

▶ Global Sum — Compute $s = \sum_{i=1}^{n} a_i$, $1 \le i \le n$

```
s \leftarrow a_1
for i = 2 to n do
    s \leftarrow s + a_i
```

Total 3n time units:

- n reads (for a_1, \ldots, a_n);
- n computes (for s); - n writes (for s).
- ▶ Prefix Sums Compute $s_i = \sum_{i=1}^k a_i, \ 1 \le i \le n$

$$\begin{aligned} s_1 &\leftarrow a_1 \\ \text{for } i &= 2 \text{ to } n \text{ do} \\ s_i &\leftarrow s_{i-1} + a_i \\ \text{endfor} \end{aligned}$$

Total 3n time units:

- n reads (for a_1, \ldots, a_n);
- n computes (for s_1, \ldots, s_n);
- n writes (for s_1, \ldots, s_n).

Jingke Li (Portland State University) CS 415/515 PRAM and Other Models

5 / 24

The Parallel RAM (PRAM) Model

- ► Machine Model:
 - ▶ A number of identical processors, P_1, P_2, \dots, P_N .
 - A common memory with M cells, shared by the N processors.
 - ▶ The processors work in a synchronous fashion.
- Basic Operations:
 - ▶ Read (Up to N) processors read simultaneously from memory. Each reads from one memory cell and stores the value in a local register.
 - ► Compute (Up to N) processors perform an (arithmetic) operation on their local data in register(s).
 - ► Write (Up to N) processors write simultaneously into register into an arbitrary memory cell.

Inter-processor communication is not explicitly modeled.

- Cost Model:
 - ▶ The compute operation takes one time unit to execute.
 - ▶ The reads and writes each takes one time unit to execute when there is no conflicts. (Otherwise, more info is needed.)

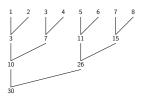
Jingke Li (Portland State University) CS 415/515 PRAM and Other Models

PRAM Algorithm Examples

Complexity is measured by $time \times space$ (#processors).

► Global Sum — final result in s₁

```
\overline{\operatorname{\mathsf{spawn}}(P_1, P_2, \dots, P_{\lceil n/2 \rceil})}
for all P_i do
       s_i \leftarrow a_i
       for j \leftarrow 1 to \lceil \log n \rceil do
             if i \% 2^j = 1 and i + 2^{j-1} \le n
                    s_i \leftarrow s_i + s_{i+2i-1}
             endif
       endfor
endfor
```

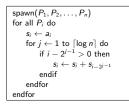


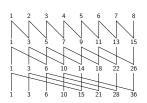
No read or write conflicts, so Time = $\lceil \log n \rceil$, Space = $\lceil \frac{n}{2} \rceil$

Jingke Li (Portland State University) CS 415/515 PRAM and Other Models

PRAM Algorithm Examples (cont.)

► Prefix Sum:





There is no write conflict, but there are potential read conflicts. Assume it's OK for now. Then, Time = $\lceil \log n \rceil$, Space = n

Jingke Li (Portland State University) CS 415/515 PRAM and Other Models

6 / 24

Cost-Optimal PRAM Algorithms

How can we compare performance between a parallel algorithm and a sequential algorithm? Or between two parallel algorithms?

The cost-optimal metric can help:

- ▶ Cost The cost of a PRAM algorithm is the product of its time complexity and space complexity.
- ► Cost-Optimal A cost-optimal PRAM algorithm is one in which the cost is in the same complexity class as the optimal sequential algorithm.

Jingke Li (Portland State University)

CS 415/515 PRAM and Other Mode

Optimality Analysis

- ► Global Sum (PRAM):

 - Time = $\lceil \log n \rceil$, Space = $\lceil \frac{n}{2} \rceil$ Cost = $\lceil \log n \rceil \times \lceil \frac{n}{2} \rceil = O(n \log n)$ Sequential lower bound = $\Theta(n)$

Conclusion: Not optimal!

Reason: Some processors have idling steps — the total number of operations performed is n-1. How to improve?

- Prefix Sums (PRAM):
 - Time = $\lceil \log n \rceil$, Space = n
 - Cost = $\lceil \log n \rceil \times n = O(n \log n)$
 - Sequential lower bound = $\Theta(n)$

Conclusion: Not optimal!

But there is almost no idle steps — the total number of operations performed is $n\lceil \log n \rceil$! How to explain?

Jingke Li (Portland State University) CS 415/515 PRAM and Other Models

Brent's Theorem

Key Idea: Any PRAM algorithm can be simulated by a different PRAM with fewer processors.

Theorem. Let A be a PRAM algorithm with time complexity t. If A performs m total operations, then p processors can simulate the PRAM and execute A in time t + (m - t)/p.

Proof:

- Let s_i be the number of operations performed by A at step i, i.e. i varies from 1 to t.
- ▶ By definition, $\sum_{i=1}^{t} s_i = m$.
- ▶ Using p processors, we can simulate step i in time $\lceil s_i/p \rceil$.
- ▶ The entire computation can then be done in time

$$\sum_{i=1}^t \lceil \frac{s_i}{p} \rceil \leq \sum_{i=1}^t \frac{s_i - 1 + p}{p} = \sum_{i=1}^t \frac{p}{p} + \sum_{i=1}^t \frac{s_i - 1}{p} = t + \frac{m - t}{p}$$

Jingke Li (Portland State University) CS 415/515 PRAM and Other Models

11 / 24

Applying Brent's Theorem

Global Sum:

- ▶ Choose a new PRAM with $p = \lfloor n/\log n \rfloor$ processors
- ▶ Total number of operations performed is n-1

$$\lceil \log n \rceil + \frac{n-1 - \lceil \log n \rceil}{\lfloor n/\log n \rfloor} = \Theta(2\log n - \frac{\log n}{n} - \frac{\log^2 n}{n}) = \Theta(\log n)$$

▶ Cost = $|n/\log n| \times \Theta(\log n) = \Theta(n)$ — It is optimal!

Jingke Li (Portland State University) CS 415/515 PRAM and Other Models

Applying Brent's Theorem (cont.)

Prefix Sums (PRAM):

- ▶ Choose a new PRAM with $p = |n/\log n|$ processors
- ▶ Total number of operations performed is $n \lceil \log n \rceil$
- ► The new time is

$$\lceil \log n \rceil + \frac{n \lceil \log n \rceil - \lceil \log n \rceil}{\lfloor n / \log n \rfloor} = \Theta(\log n + \log^2 n - \frac{\log^2 n}{n}) = \Theta(\log^2 n)$$

▶ Cost = $\lfloor n/\log n \rfloor \times \Theta(\log^2 n) = \Theta(n \log n)$ — Still not optimal!

Reason: The total number of PRAM operations is higher than that of the optimal sequential algorithm.

"Coarse-Grain" Algorithms

Another approach for improving performance.

Prefix Sums (Coarse-Grain) — use < n processors

- Divide the *n* values into *p* sets, each containing $\leq \lceil n/p \rceil$ values.
- The first $\it p-1$ processors each uses the optimal sequential algorithm to do a local prefix computation. (time: $\lceil n/p \rceil - 1$)
- Then the processors run the PRAM parallel prefix algorithm on the subtotals. (time: log(p-1))
- Each processor then goes back and updates the local prefix values with the results from the global comp. (time: $\lceil n/p \rceil$)

Total Cost: $(2\lceil n/p \rceil + \log(p-1))p = \Theta(n+p\log p)$

For small p, this algorithm is cost-optimal!

Jingke Li (Portland State University) CS 415/515 PRAM and Other Models

Jingke Li (Portland State University) CS 415/515 PRAM and Other Models

PRAM Memory Access Models

- ► Exclusive Read (ER)
 - For any memory cell, only a single read is allowed at any given time.
- ► Concurrent Read (CR)
 - Simultaneous multiple reads to the same memory cell are allowed.
- ► Exclusive Write (EW)
 - For any memory cell, only a single write is allowed at any given time.
- ► Concurrent Write (CW)
 - Simultaneous multiple writes to the same memory cell are allowed.
 - Question: What value gets written?

Jingke Li (Portland State University) CS 415/515 PRAM and Other Models

13 / 24

Concurrent Write Sub Models

- ► Priority CW
 - The processors are assigned certain priorities; the processor with the highest priority wins.
- Common CW
 - The processors are allowed to write to the same memory cell only if they are attempting to write the same value; otherwise a special flag will be raised.
- Arbitrary CW
 - Any of the attempting processors can succeed; the selection is according to an algorithm.
- Random CW
 - A processor is chosen by a random process to succeed.
- Combining CW
 - All the values from the attempting processors are combined into a single value, which is then stored in the memory cell.
 - Combing operators: sum, product, and, max, min, etc.

Jingke Li (Portland State University)

CS 415/515 PRAM and Other Models

Three Specific PRAM Models

These three models are commonly used in parallel algorithm analysis:

- ▶ EREW PRAM The weakest among the three models. It takes $O(\log n)$ time units to spread a value from one processor to n other processors
- ► CREW PRAM A processor can spread a value to n other processors in O(1) time units.
- ► CRCW PRAM The strongest model. But CW needs special handling for resolving memory write conflicts.

Jingke Li (Portland State University)

CS 415/515 PRAM and Other Models

15 / 24

17 / 24

Concurrent-Write Algorithms

► Global Sum (CR+Combing CW PRAM):

Time = O(1)Space = n

▶ Prefix Sums (CR+Combining CW PRAM):

$$\begin{array}{l} \operatorname{spawn}(P_{i,j}:1\leq i\leq n,1\leq j\leq n-i)\\ \text{for all }P_{i,j}\text{ do}\\ t_{i,j}=a_j;\\ s_i\leftarrow t_{i,j}\\ \text{endfor} \end{array}$$

Time = O(1) $Space = n^2$

Jingke Li (Portland State University)

Critiques on the PRAM Model

zero latency, and zero overhead.

predicting the performance of algorithms.

any cell in unit time.

CS 415/515 PRAM and Other Models

It assumes the processors operate in a fully-synchronous mode.

▶ It neglects the issue of contention caused by concurrent access to

▶ It assumes the interprocessor communication has infinite bandwidth,

▶ It assumes the number of processors can increase with the problem

In conclusion, the PRAM model is good for a gross classification of

algorithms, but not very useful for describing realistic algorithms or

different cells within the same memory module.

▶ It assumes a single shared memory in which each processor can access

Two Rank Sort Examples

► Rank Sort (CREW PRAM):

```
spawn(P_1, P_2, \dots, P_n)
    k = 0
    for j=0 to n do
        if a[i] > a[j] then k \leftarrow k+1
    b[k] \leftarrow a[i]
```

Time = O(n)Space = n

► Rank Sort (CR+Combining CW PRAM):

$$\begin{aligned} & \mathsf{spawn}(P_1, P_2, \dots, P_n) \\ & \mathsf{for all} \ P_i \ \mathsf{do} \\ & k = 0; \ \mathsf{spawn}(Q_1, Q_2, \dots, Q_n) \\ & \mathsf{for all} \ j = 1 \ \mathsf{to} \ n \ \mathsf{do} \\ & \mathsf{if} \ a[i] > a[j] \ \mathsf{then} \ k \leftarrow 1 \\ & b[k] \leftarrow a[i] \\ & \mathsf{endfor} \end{aligned}$$

Time = O(1) $Space = n^2$

Jingke Li (Portland State University)

CS 415/515 PRAM and Other Models

Jingke Li (Portland State University) CS 415/515 PRAM and Other Models

18 / 24

Extensions of the PRAM Model

- ▶ Phase PRAM Introduces asynchrony.
- ▶ Module Parallel Computer Divides memory into modules.
- ► Local-Memory PRAM Divides memory into local and global.
- ▶ Memory Hierarchy Model Views the memory as a hierarchy.
- ► Delay Model Introduces communication delay.

Jingke Li (Portland State University) CS 415/515 PRAM and Other Models

19 / 24

Valiant's BSP Model

BSP = Bulk Synchronous Parallel

The BSP model is developed to model message-passing multicomputers. It consists of three components:

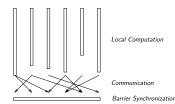
- A group of p processors each with local memory.
- ► An interconnection network for *point-to-point communication* between the processors.
- A mechanism for synchronizing all the processors at defined intervals.

Jingke Li (Portland State University) CS 415/515 PRAM and Other Models

Supersteps

In BSP, a computation consists of a sequence of supersteps. Each superstep is further subdivided into three ordered phases:

- ► Local Computation each processor performs computation using only data stored in the local memory.
- ► Communication processors sends/receives messages to each other.



▶ Barrier Synchronization — this global synchronization waits for all of the communication actions to complete.

Jingke Li (Portland State University) CS 415/515 PRAM and Other Models

21 / 24

BSP Architectural Parameters

BSP model has three architectural parameters:

p — the number of processors

I — the latency of a barrier synchronization

g — the communication overhead

They can be computed for a particular real target machine.

With these parameters, the time of a superstep can be expressed by

$$T_{superstep} = w_{max} + gh_{max} + I$$

where w_{max} is the maximum time caused by local operations by any processor and h_{max} is the maximum number of messages being sent or received by a processor.

Jingke Li (Portland State University) CS 415/515 PRAM and Other Models

BSP Model Properties

- ► High-level architecture-independent.
- ▶ Simple to use BSP programs look much the same as sequential programs.
- ▶ Can predict performance With a small set of parameters used to describe a given architecture, the performance of a BSP program on the target architecture is predictable.

Jingke Li (Portland State University) CS 415/515 PRAM and Other Models

23 / 24

A BSP Program

The following function calculates the partial sums of p integers stored on pprocessors. (The code is written in C with Oxford BSP library.)

```
int bsp_allsums(int x) {
 nt bsp_allsums(int x) {
int i, left, right;
int mypid = bsp_pid();
int p = bsp_nprocs();
bsp_pushregister(&left, sizeof(int));
bsp_sync();
  right = x;
for (i=1; i<p; i*=2) {
     if (mypid+i < p)
bsp_put(mypid+i, &right, &left, 0, sizeof(int));</pre>
     bsp_sync();
if (mypid>=i)
        right = left + right;
  bsp_popregister(&left);
  return right;
```

Jingke Li (Portland State University) CS 415/515 PRAM and Other Models

24 / 24