

## Numerical Algorithms

Jingke Li

Portland State University

## Basic Numerical Algorithms

- *Dense-Matrix Algorithms:*
  - Matrix-vector multiplication
  - Matrix-matrix multiplication
  - Gaussian elimination
- *Sparse-Matrix Algorithms:*
  - Jacobi relaxation
  - Gauss-Seidel
  - Multi-grid method

For all these algorithms, parallelizing them to run on a shared-memory system is trivial. They all have easily-parallelizable loops.

## Matrix-Vector Multiplication

Compute  $y = Ax$  ( $A$  is an  $n \times n$  matrix,  $x, y$  are  $n \times 1$  vectors)

*Sequential Algorithm:*

```
void matrix_vector(int n, double a[n][n], double x[n], double y[n]) {
    int i, j;
    for (i = 0; i < n; i++) {
        y[i] = 0;
        for (j = 0; j < n; j++)
            y[i] += a[i][j] * x[j];
    }
}
```

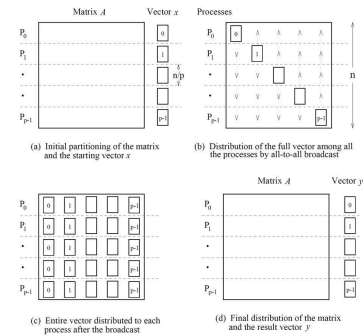
*Complexity:*  $T = O(n^2)$

*Parallelization (for Distributed-Memory Systems):*

- Need to decide partition and communication.

## Distributed-Memory Matrix-Vector Multiplication

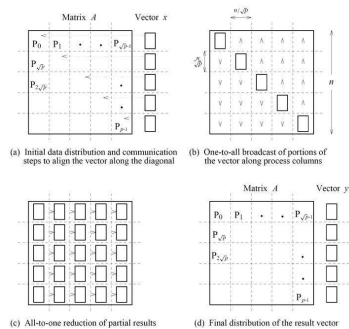
- 1D Partition (row-wise):



*Complexity:*  $T_p \approx \frac{n^2}{p} + t_s \log p + t_w n$

## Distributed-Memory Matrix-Vector Multiplication (cont.)

- 2D Partition:



*Complexity:*  $T_p \approx \frac{n^2}{p} + t_s \log p + t_w \frac{n}{\sqrt{p}} \log p$

## Matrix-Matrix Multiplication

Compute  $C = AB$  ( $A, B, C$  are  $n \times n$  dense matrices)


*Sequential Algorithm:*

```
void matrix_mult(int n, double a[n][n], double b[n][n], double c[n][n]) {
    int i, j, k;
    for (i = 0; i < n; i++) {
        for (j = 0; j < n; j++) {
            c[i][j] = 0;
            for (k = 0; k < n; k++)
                c[i][j] += a[i][k] * b[k][j];
        }
    }
}
```

*Complexity:*  $T = O(n^3)$  (Best is  $O(n^{2.8})$  — Strassen's Algorithm)

## Distributed-Memory Matrix-Matrix Multiplication

### Partition Choices:



### Communication Choices:

- A single broadcast — The needed data is broadcast to all destinations once for all.
  - A single step, but needs a lot of buffer storage.
- Multiple shifts — The needed data is shifted towards its destination one step at a time.
  - Only local communication and very little buffer space, but takes multiple steps for data to reach destination.


## A Simple 2D Matrix-Mult Algorithm

- Processors are arranged in a logical  $\sqrt{p} \times \sqrt{p}$  2D topology.

- Each processor gets a block of  $\frac{n}{\sqrt{p}} \times \frac{n}{\sqrt{p}}$  block of A, B, & C.

- Perform all-to-all broadcasts within processor rows for A's blocks.

- Take  $\sqrt{p}$  iterations of the following steps:
  1. Multiply A's block with B's block to form C's block.
  2. Shift B's blocks within processor columns.



$$\text{Complexity: } T_p \approx \frac{n^3}{p} + t_s \log p + 2t_w \frac{n^2}{\sqrt{p}}$$

## Cannon's Matrix-Mult Algorithm

Memory efficient variant of the simple algorithm.

Key Idea: Replace the standard loop

$$C_{i,j} = \sum_{k=0}^{\sqrt{p}-1} A_{i,k} \times B_{k,j}$$

with a skewed loop


$$C_{i,j} = \sum_{k=0}^{\sqrt{p}-1} A_{i,(i+j+k) \bmod \sqrt{p}} \times B_{(i+j+k) \bmod \sqrt{p},j}$$

Communication: Multiple shifts for both arrays A and B.

$$\text{Complexity: } T_p \approx \frac{n^3}{p} + 2t_s \sqrt{p} + 2t_w \frac{n^2}{\sqrt{p}}$$

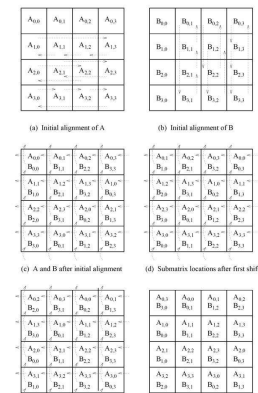
## Cannon's Matrix-Mult Algorithm (cont.)

Initial Skewing of A and B:



Iterations:

- Multiply A and B blocks.
- Shift A blocks towards left.
- Shift B blocks upwards.



## Gaussian Elimination

A well-known algorithm for solving a linear system  $\mathbf{Ax} = \mathbf{b}$ .

Idea:

- First reduce  $\mathbf{Ax} = \mathbf{b}$  to an upper triangular system  $\mathbf{T}\mathbf{x} = \mathbf{c}$ .
- Then use back substitution to solve for  $\mathbf{x}$ .

The two needed operations are:

- Interchange any two rows (this simply reorders the  $n$  equations).
- Replace any row by a linear combination of itself and another row.

## Gaussian Elimination Example

- The initial linear system:

$$\begin{aligned} 2.0 x_1 + 1.0 x_2 + 3.0 x_3 + 4.0 x_4 &= 29.0 \\ 5.0 x_1 + 2.0 x_2 + 0.0 x_3 + 1.0 x_4 &= 13.0 \\ 1.0 x_1 + 0.0 x_2 + 0.0 x_3 + 1.0 x_4 &= 5.0 \\ 3.0 x_1 + 1.0 x_2 + 1.0 x_3 + 0.0 x_4 &= 8.0 \end{aligned}$$

- Matrix form and triangulating:

$$\begin{bmatrix} 2.0 & 1.0 & 3.0 & 4.0 & 29.0 \\ 5.0 & 2.0 & 0.0 & 1.0 & 13.0 \\ 1.0 & 0.0 & 0.0 & 1.0 & 5.0 \\ 3.0 & 1.0 & 1.0 & 0.0 & 8.0 \end{bmatrix} \Rightarrow \begin{bmatrix} 2.0 & 1.0 & 3.0 & 4.0 & 29.0 \\ 0.0 & -0.5 & -7.5 & -9.0 & -59.5 \\ 0.0 & 0.0 & 6.0 & 8.0 & 50.0 \\ 0.0 & 0.0 & 0.0 & -2.3 & -9.3 \end{bmatrix}$$

- Back to linear system form:

$$\begin{aligned} 2.0 x_1 + 1.0 x_2 + 3.0 x_3 + 4.0 x_4 &= 29.0 \\ -0.5 x_2 - 7.5 x_3 - 9.0 x_4 &= -59.5 \\ 6.0 x_3 + 8.0 x_4 &= 50.0 \\ -2.3 x_4 &= -9.3 \end{aligned}$$


- Use back-substitution to solve:  $x_1=1, x_2=2, x_3=3, x_4=4$

## Naive Gaussian Elimination

```
void gaussian_naive(int n, double a[n][n], double b[n], double x[n]) {
    int i, j, k;
    for (k = 0; k < n; k++) {
        for (j = k+1; j < n; j++) {
            a[k][j] = a[k][j] / a[k][k];
        }
        x[k] = b[k] / a[k][k];
        for (i = k+1; i < n; i++) {
            for (j = k+1; j < n; j++) {
                a[i][j] -= a[i][k] * a[k][j];
            }
            b[i] = a[i][k] * x[k];
            a[i][k] = 0;
        }
    }
}
```

Complexity:  $T = O(n^3)$

## Naive Gaussian Elimination (cont.)



## Gaussian Elimination with Partial-Pivoting

The naive GE algorithm runs into trouble when a pivot element is zero. (In fact, it runs into trouble whenever a pivot element is “small.”)

**Partial-Pivoting** — Search the column below the diagonal and find the largest element in absolute magnitude; perform row interchange to bring this element to the diagonal.

**Example:**

Initial:                      => Pivoting:

2.0	1.0	3.0	4.0	29.0	5.0	2.0	0.0	1.0	13.0
1.0	0.0	0.0	1.0	5.0	1.0	0.0	0.0	1.0	5.0
3.0	1.0	1.0	0.0	8.0	3.0	1.0	1.0	0.0	8.0
5.0	2.0	0.0	1.0	13.0	2.0	1.0	3.0	4.0	29.0

=> Elimination:                      =>

5.0	2.0	0.0	1.0	13.0					
0.0	-0.4	0.0	0.8	2.4					
0.0	-0.2	1.0	-0.6	0.2					
0.0	0.2	3.0	3.6	23.8					

Continue for other columns ...

## Parallel GE with 1-D Partitioning

$P_0$	1	(0.1)	(0.2)	(0.3)	(0.4)	(0.5)	(0.6)	(0.7)
$P_1$	0	1	(1.2)	(1.3)	(1.4)	(1.5)	(1.6)	(1.7)
$P_2$	0	0	1	(2.3)	(2.4)	(2.5)	(2.6)	(2.7)
$P_3$	0	0	0	(3.3)	(3.4)	(3.5)	(3.6)	(3.7)
$P_4$	0	0	0	0	(4.4)	(4.5)	(4.6)	(4.7)
$P_5$	0	0	0	0	0	(5.4)	(5.5)	(5.6)
$P_6$	0	0	0	0	0	0	(6.6)	(6.7)
$P_7$	0	0	0	0	0	0	0	(7.7)

- (a) Computation:
- (i)  $A[k,j] := A[k,j]/A[k,k]$  for  $k < j < n$
  - (ii)  $A[k,k] := 1$

- (b) Communication:
- One-to-all broadcast of row  $A[k,*]$

- (c) Computation:
- (i)  $A[i,j] := A[i,j] - A[i,k] * A[k,j]$  for  $k < i < n$  and  $k < j < n$
  - (ii)  $A[i,k] := 0$  for  $k < i < n$

## Parallel GE with 1-D Partitioning

- Assign one row to each process.
- Execute the outer loop sequentially.
- At iteration  $k+1$ , process  $P_k$  either broadcast or shift its row to processes  $P_{k+1}, \dots, P_{n-1}$ .
- Each process performs local computation.

Complexity:

$$T_p = \frac{3}{2}n(n-1) + t_s n \log n + \frac{1}{2}t_w n(n-1) \log n$$

## Improvements

**Problems with the Simple Algorithm** — The partition is fine-grained, and towards the end, the active region of the matrix is shrinking towards lower right corner; implies that processes fall idle.

**Solutions:** block and cyclic partitions.

$P_0$	1	(0.1)	(0.2)	(0.3)	(0.4)	(0.5)	(0.6)	(0.7)
$P_1$	0	1	(1.2)	(1.3)	(1.4)	(1.5)	(1.6)	(1.7)
$P_2$	0	0	1	(2.3)	(2.4)	(2.5)	(2.6)	(2.7)
$P_3$	0	0	0	1	(3.4)	(3.5)	(3.6)	(3.7)
$P_4$	0	0	0	0	1	(4.5)	(4.6)	(4.7)
$P_5$	0	0	0	0	0	1	(5.6)	(5.7)
$P_6$	0	0	0	0	0	0	1	(6.7)
$P_7$	0	0	0	0	0	0	0	1

(a) Block 1-D mapping

$P_0$	1	(0.1)	(0.2)	(0.3)	(0.4)	(0.5)	(0.6)	(0.7)
$P_1$	0	0	0	(4.3)	(4.4)	(4.5)	(4.6)	(4.7)
$P_2$	0	1	(1.2)	(1.3)	(1.4)	(1.5)	(1.6)	(1.7)
$P_3$	0	0	0	(5.3)	(5.4)	(5.5)	(5.6)	(5.7)
$P_4$	0	0	1	(2.3)	(2.4)	(2.5)	(2.6)	(2.7)
$P_5$	0	0	0	(6.3)	(6.4)	(6.5)	(6.6)	(6.7)
$P_6$	0	0	0	(3.3)	(3.4)	(3.5)	(3.6)	(3.7)
$P_7$	0	0	0	(7.3)	(7.4)	(7.5)	(7.6)	(7.7)

(b) Cyclic 1-D mapping

## Parallel GE with Partial-Pivoting

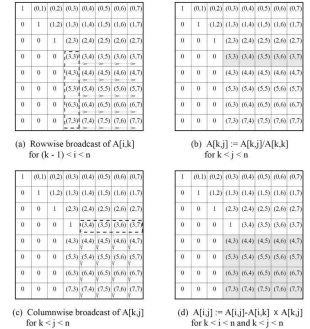
- Distribute the matrix as complete rows; each process stores approximately  $n/p$  rows of the matrix.
- Processes collectively decide on which two rows need to be swapped to do the partial pivoting.
- Two processes do a `send/recv` to each other to do the swap. (Alternatively, we can keep track of which row is which through an indirection array.)
- The pivot row is broadcast to all processes.
- In each process, we loop over those rows below the pivot row and apply the elimination step (an `saxby()` operation).
- Repeat until we hit bottom.

## Parallel GE with 2-D Partitioning

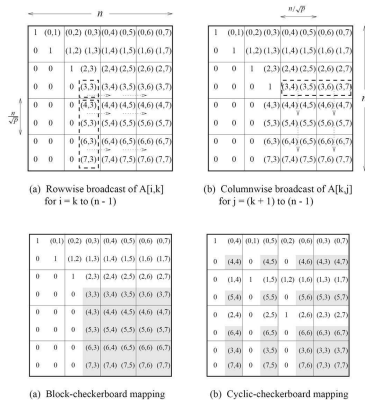
Each processor gets a 2D block of the matrix.

Steps:

- Broadcast the “active” column along the rows.
- Prepare the pivot row concurrently.
- Broadcast the pivot row along the columns.
- Perform the elimination step concurrently.



## Block and Cyclic 2-D Partitions




## Solving Large, Sparse Linear Systems

Large sparse matrices often appear in scientific or engineering applications when solving partial differential equations. As an example, consider the Laplace Equation:

$$\phi_{i,j} = \frac{1}{4}(\phi_{i+1,j} + \phi_{i-1,j} + \phi_{i,j+1} + \phi_{i,j-1}), \quad 0 < i, j < n$$

It describes a computation over an  $n \times n$  mesh.  $\phi_{i,j}$  denotes the value at the mesh point  $[i, j]$ .



- The value at point  $[i, j]$  is related to the values of its four neighbors.
- The values at the four boarder lines are typically fixed.

## Solving the Laplace Equation

The Laplace Equation is a linear system with  $n^2$  unknowns. It can be turned into the standard  $\mathbf{Ax} = \mathbf{b}$  form:

Let  $l = i \cdot n + j$ , then the system becomes:

$$\phi_{l+n} + \phi_{l-n} + \phi_{l+1} + \phi_{l-1} - 4\phi_l = 0, \quad 0 < l < n^2$$

which can also be expressed in matrix form:

$$\begin{bmatrix} -4 & 1 & 0 & 0 & \dots & 1 & 0 & \dots \\ 1 & -4 & 1 & 0 & \dots & 0 & 1 & \dots \\ 0 & 1 & -4 & 1 & \dots & 0 & 0 & \dots \\ 0 & 0 & 1 & -4 & \dots & \dots & \dots & \dots \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \vdots \\ 1 & 0 & 0 & \dots & \dots & \dots & \dots & \dots \\ 0 & 1 & 0 & \dots & \dots & \dots & \dots & \dots \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \vdots \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \end{bmatrix} \mathbf{x} = \mathbf{0}$$

- The matrix is very large, yet sparse.
- It can be solved by using Gaussian Elimination, but the cost would be very high:  $O(n^6)$ .

A better approach is to directly solve the system over the original  $n \times n$  mesh domain.

## Jacobi Relaxation Algorithm

It is an iterative algorithm:

- Cycle through the mesh points, compute a new value for each point using the average of the four-neighboring “old” values.

$$\phi_{i,j}^{(t+1)} = \frac{1}{4}(\phi_{i+1,j}^{(t)} + \phi_{i-1,j}^{(t)} + \phi_{i,j+1}^{(t)} + \phi_{i,j-1}^{(t)}), \quad 0 < i, j < N$$

- Once all the new values are found, replace old values with new ones.
- Repeat until the difference between old and new is small enough.

Parallelization:

- For the distributed-memory case, simply partition the mesh in both dimensions; each processor will handle a  $\frac{n}{\sqrt{p}} \times \frac{n}{\sqrt{p}}$  sub-mesh. Communications are only required on the “peripheral” of the mesh region in each processor.

## Jacobi Relaxation Algorithm (cont.)

```
int jacobi(int n, double x[n][n], double epsilon) {
    double xnew[n][n], delta;
    int i, j, cnt = 0;
    do {
        cnt++;
        delta = 0.0;
        for (i = 1; i < n-1; i++) {
            for (j = 1; j < n-1; j++) {
                xnew[i][j] = (x[i-1][j] + x[i][j-1]
                    + x[i+1][j] + x[i][j+1]) / 4.0;
                delta = fmax(delta, fabs(xnew[i][j] - x[i][j]));
            }
        }
        for (i = 1; i < n-1; i++) {
            for (j = 1; j < n-1; j++) {
                x[i][j] = xnew[i][j];
            }
        }
    } while (delta > epsilon);
    return cnt;
}
```

## Gauss-Seidel Algorithm

While simple, the Jacobi relaxation algorithm has two drawbacks, (1) It needs buffers to store old mesh point values; and (2) Its convergence speed can be slow.

Gauss-Seidel algorithm is an improvement over the Jacobi algorithm. It's like Jacobi, except that mesh points are updated "in-place," overwriting the old value.

The order of mesh point updates does not really matter. Here are two possible orders:

$$\phi_{i,j}^{(t+1)} = \frac{1}{4}(\phi_{i+1,j}^{(t)} + \phi_{i-1,j}^{(t+1)} + \phi_{i,j+1}^{(t)} + \phi_{i,j-1}^{(t+1)}), \quad 0 < i, j < N$$

$$\phi_{i,j}^{(t+1)} = \frac{1}{4}(\phi_{i+1,j}^{(t+1)} + \phi_{i-1,j}^{(t)} + \phi_{i,j+1}^{(t+1)} + \phi_{i,j-1}^{(t)}), \quad 0 < i, j < N$$

- Gauss-Seidel algorithm does not need buffers and has a better convergence speed (since newer values are used in all updates).

## Parallelizing Gauss-Seidel

Due to the "in-place" updates, race condition may happen among multiple reads and a single write to the same memory location. But this is not really a problem for an iterative algorithm — using an old value vs. using a new value affects only the convergence rate.

The following are the common parallel approaches:

- *Natural index ordering* (allow race condition)

- *Red-black ordering* (avoid race condition)

- Denote alternating points as "red" and "black".
- Repeat following two steps until convergence:
  - update all red points (in any order)
  - update all black points (in any order)

```
r b r b r b
b r b r b r
r b r b r b
b r b r b r
r b r b r b
b r b r b r
```

- *Wavefront ordering* (avoid race condition)

- Proceed along diagonal line.
- Repeat until success.

```
1 2 3 4 ..
2 3 4 5 ..
3 4 5 ..
4 5 ..
5 ..
..
```

## Multi-Grid Methods

*Observations* — Iterative algorithms converge to solutions faster on coarser grids than on finer grids. Iterative algorithms converge quicker if the initial approx. of the values of the variables are good.

*Algorithm:*

- Begin with the original problem, where the solution is defined (and desired) on an  $n \times n$  mesh.
- (*Projection*) Coarsen the problem by several powers of 2:  $n \rightarrow n/2^m$ . Boundary values need to be averaged down to the coarser mesh.
- (*Relaxation*) Solve the coarse version problem using any iterative solver. It should go fast since the mesh is small.
- (*Interpolation*) Boost up the problem by a factor of 2, interpolating field points. What was one mesh point turns into 4 mesh points.
- Re-run relaxation on this problem.
- Repeat *Relaxation/Interpolation* steps until back to original problem.