

Numerical Algorithms

Jingke Li

Portland State University

Basic Numerical Algorithms

- ▶ *Dense-Matrix Algorithms:*
 - ▶ Matrix-vector multiplication
 - ▶ Matrix-matrix multiplication
 - ▶ Gaussian elimination
- ▶ *Sparse-Matrix Algorithms:*
 - ▶ Jacobi relaxation
 - ▶ Gauss-Seidel
 - ▶ Multi-grid method

For all these algorithms, parallelizing them to run on a shared-memory system is trivial. They all have easily-parallelizable loops.

Matrix-Vector Multiplication

Compute $y = Ax$ (A is an $n \times n$ matrix, x, y are $n \times 1$ vectors)

Sequential Algorithm:

```
void matrix_vector(int n, double a[n][n], double x[n], double y[n]) {
    int i, j;
    for (i = 0; i < n; i++) {
        y[i] = 0;
        for (j = 0; j < n; j++)
            y[i] += a[i][j] * x[j];
    }
}
```

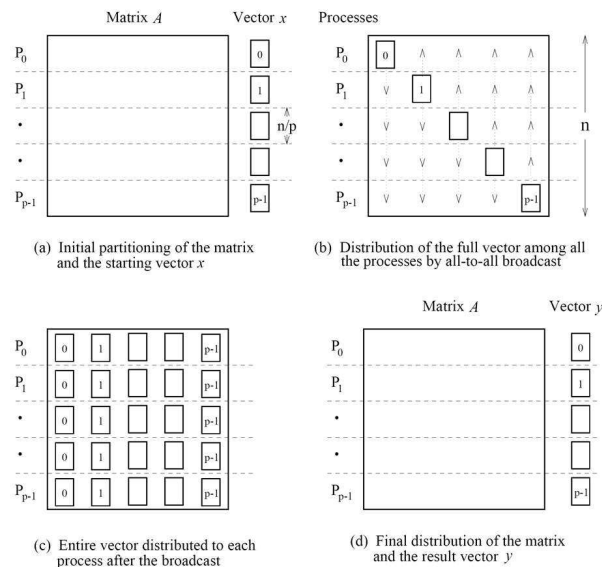
Complexity: $T = O(n^2)$

Parallelization (for Distributed-Memory Systems):

- Need to decide partition and communication.

Distributed-Memory Matrix-Vector Multiplication

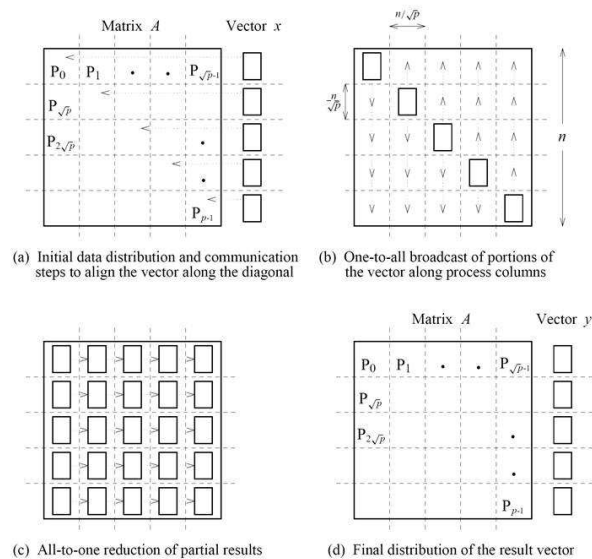
- 1D Partition (row-wise):



Complexity: $T_p \approx \frac{n^2}{p} + t_s \log p + t_w n$

Distributed-Memory Matrix-Vector Multiplication (cont.)

► 2D Partition:



$$\text{Complexity: } T_p \approx \frac{n^2}{p} + t_s \log p + t_w \frac{n}{\sqrt{p}} \log p$$

Matrix-Matrix Multiplication

Compute $C = AB$ (A, B, C are $n \times n$ dense matrices)


Sequential Algorithm:

```
void matrix_mult(int n, double a[n][n], double b[n][n], double c[n][n]) {
    int i, j, k;
    for (i = 0; i < n; i++) {
        for (j = 0; j < n; j++) {
            c[i][j] = 0;
            for (k = 0; k < n; k++)
                c[i][j] += a[i][k] * b[k][j];
        }
    }
}
```

Complexity: $T = O(n^3)$ (Best is $O(n^{2.8})$ — Strassen's Algorithm)

Distributed-Memory Matrix-Matrix Multiplication

► Partition Choices:

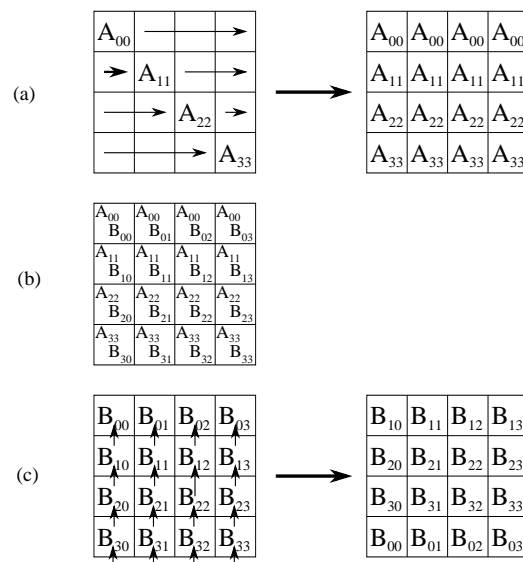


► Communication Choices:

- A single broadcast — The needed data is broadcast to all destinations once for all.
 - A single step, but needs a lot of buffer storage.
- Multiple shifts — The needed data is shifted towards its destination one step at a time.
 - Only local communication and very little buffer space, but takes multiple steps for data to reach destination.

A Simple 2D Matrix-Mult Algorithm

- Processors are arranged in a logical $\sqrt{p} \times \sqrt{p}$ 2D topology.
- Each processor gets a block of $\frac{n}{\sqrt{p}} \times \frac{n}{\sqrt{p}}$ block of A , B , & C .
- Perform all-to-all broadcasts within processor rows for A 's blocks.
- Take \sqrt{p} iterations of the following steps:
 1. Multiply A 's block with B 's block to form C 's block.
 2. Shift B 's blocks within processor columns.



Complexity: $T_p \approx \frac{n^3}{p} + t_s \log p + 2t_w \frac{n^2}{\sqrt{p}}$

Cannon's Matrix-Mult Algorithm

Memory efficient variant of the simple algorithm.

Key Idea: Replace the standard loop

$$C_{i,j} = \sum_{k=0}^{\sqrt{p}-1} A_{i,k} \times B_{k,j}$$

with a skewed loop

$$C_{i,j} = \sum_{k=0}^{\sqrt{p}-1} A_{i,(i+j+k) \bmod \sqrt{p}} \times B_{(i+j+k) \bmod \sqrt{p},j}$$

Communication: Multiple shifts for both arrays A and B .

Complexity: $T_p \approx \frac{n^3}{p} + 2t_s\sqrt{p} + 2t_w\frac{n^2}{\sqrt{p}}$

Cannon's Matrix-Mult Algorithm (cont.)

Initial Skewing of A and B:

A_{00}	A_{01}	A_{02}	A_{03}
A_{11}	A_{12}	A_{13}	A_{10}
A_{22}	A_{23}	A_{20}	A_{21}
A_{33}	A_{30}	A_{31}	A_{32}

B_{00}	B_{11}	B_{22}	B_{33}
B_{10}	B_{21}	B_{32}	B_{03}
B_{20}	B_{31}	B_{02}	B_{13}
B_{30}	B_{01}	B_{12}	B_{23}

Iterations:

- Multiply A and B blocks.
- Shift A blocks towards left.
- Shift B blocks upwards.

$A_{0,0}$	$A_{0,1}$	$A_{0,2}$	$A_{0,3}$
$A_{1,0}$	$A_{1,1}$	$A_{1,2}$	$A_{1,3}$
$A_{2,0}$	$A_{2,1}$	$A_{2,2}$	$A_{2,3}$
$A_{3,0}$	$A_{3,1}$	$A_{3,2}$	$A_{3,3}$

(a) Initial alignment of A

$B_{0,0}$	$B_{0,1}$	$B_{0,2}$	$B_{0,3}$
$B_{1,0}$	$B_{1,1}$	$B_{1,2}$	$B_{1,3}$
$B_{2,0}$	$B_{2,1}$	$B_{2,2}$	$B_{2,3}$
$B_{3,0}$	$B_{3,1}$	$B_{3,2}$	$B_{3,3}$

(b) Initial alignment of B

$A_{0,0}$	$A_{0,1}$	$A_{0,2}$	$A_{0,3}$
$B_{0,0}$	$B_{1,1}$	$B_{2,2}$	$B_{3,3}$
$A_{1,1}$	$A_{1,2}$	$A_{1,3}$	$A_{1,0}$
$B_{1,0}$	$B_{2,1}$	$B_{3,2}$	$B_{0,3}$
$A_{2,2}$	$A_{2,3}$	$A_{2,0}$	$A_{2,1}$
$B_{2,0}$	$B_{3,1}$	$B_{0,2}$	$B_{1,3}$
$A_{3,3}$	$A_{3,0}$	$A_{3,1}$	$A_{3,2}$
$B_{3,0}$	$B_{0,1}$	$B_{1,2}$	$B_{2,3}$

(c) A and B after initial alignment

$A_{0,1}$	$A_{0,2}$	$A_{0,3}$	$A_{0,0}$
$B_{1,0}$	$B_{2,1}$	$B_{3,2}$	$B_{0,3}$
$A_{1,2}$	$A_{1,3}$	$A_{1,0}$	$A_{1,1}$
$B_{2,0}$	$B_{3,1}$	$B_{0,2}$	$B_{1,3}$
$A_{2,3}$	$A_{2,0}$	$A_{2,1}$	$A_{2,2}$
$B_{3,0}$	$B_{0,1}$	$B_{1,2}$	$B_{2,3}$
$A_{3,0}$	$A_{3,1}$	$A_{3,2}$	$A_{3,3}$
$B_{0,0}$	$B_{1,1}$	$B_{2,2}$	$B_{3,3}$

(d) Submatrix locations after first shift

$A_{0,2}$	$A_{0,3}$	$A_{0,0}$	$A_{0,1}$
$B_{2,0}$	$B_{3,1}$	$B_{0,2}$	$B_{1,3}$
$A_{1,3}$	$A_{1,0}$	$A_{1,1}$	$A_{1,2}$
$B_{3,0}$	$B_{0,1}$	$B_{1,2}$	$B_{2,3}$
$A_{2,0}$	$A_{2,1}$	$A_{2,2}$	$A_{2,3}$
$B_{0,0}$	$B_{1,1}$	$B_{2,2}$	$B_{3,3}$
$A_{3,1}$	$A_{3,2}$	$A_{3,3}$	$A_{3,0}$
$B_{1,0}$	$B_{2,1}$	$B_{3,2}$	$B_{0,3}$

$A_{0,3}$	$A_{0,0}$	$A_{0,1}$	$A_{0,2}$
$B_{3,0}$	$B_{0,1}$	$B_{1,2}$	$B_{2,3}$
$A_{1,0}$	$A_{1,1}$	$A_{1,2}$	$A_{1,3}$
$B_{0,0}$	$B_{1,1}$	$B_{2,2}$	$B_{3,3}$
$A_{2,1}$	$A_{2,2}$	$A_{2,3}$	$A_{2,0}$
$B_{1,0}$	$B_{2,1}$	$B_{3,2}$	$B_{0,3}$
$A_{3,2}$	$A_{3,3}$	$A_{3,0}$	$A_{3,1}$
$B_{2,0}$	$B_{3,1}$	$B_{0,2}$	$B_{1,3}$

Gaussian Elimination

A well-known algorithm for solving a linear system $\mathbf{Ax} = \mathbf{b}$.

Idea:

- ▶ First reduce $\mathbf{Ax} = \mathbf{b}$ to an upper triangular system $\mathbf{Tx} = \mathbf{c}$.
- ▶ Then use back substitution to solve for \mathbf{x} .

The two needed operations are:

- ▶ Interchange any two rows (this simply reorders the n equations).
- ▶ Replace any row by a linear combination of itself and another row.

Gaussian Elimination Example

- ▶ The initial linear system:

$$\begin{array}{rcl} 2.0 x_1 + 1.0 x_2 + 3.0 x_3 + 4.0 x_4 & = & 29.0 \\ 5.0 x_1 + 2.0 x_2 + 0.0 x_3 + 1.0 x_4 & = & 13.0 \\ 1.0 x_1 + 0.0 x_2 + 0.0 x_3 + 1.0 x_4 & = & 5.0 \\ 3.0 x_1 + 1.0 x_2 + 1.0 x_3 + 0.0 x_4 & = & 8.0 \end{array}$$

- ▶ Matrix form and triangulating:

$$\begin{array}{ccccc} 2.0 & 1.0 & 3.0 & 4.0 & 29.0 \\ 0.0 & -0.5 & -7.5 & -9.0 & -59.5 \\ 0.0 & -0.5 & -1.5 & -1.0 & -9.5 \\ 0.0 & -0.5 & -3.5 & -6.0 & -35.5 \end{array} \Rightarrow \begin{array}{ccccc} 2.0 & 1.0 & 3.0 & 4.0 & 29.0 \\ 0.0 & -0.5 & -7.5 & -9.0 & -59.5 \\ 0.0 & 0.0 & 6.0 & 8.0 & 50.0 \\ 0.0 & 0.0 & 0.0 & -2.3 & -9.3 \end{array}$$

- ▶ Back to linear system form:

$$\begin{array}{rcl} 2.0 x_1 + 1.0 x_2 + 3.0 x_3 + 4.0 x_4 & = & 29.0 \\ -0.5 x_2 - 7.5 x_3 - 9.0 x_4 & = & -59.5 \\ 6.0 x_3 + 8.0 x_4 & = & 50.0 \\ -2.3 x_4 & = & -9.3 \end{array}$$


- ▶ Use back-substitution to solve: $x_1=1, x_2=2, x_3=3, x_4=4$

Naive Gaussian Elimination

```
void gaussian_naive(int n, double a[n][n], double b[n], double x[n]) {
    int i, j, k;
    for (k = 0; k < n; k++) {
        for (j = k+1; j < n; j++) {
            a[k][j] = a[k][j] / a[k][k];
        }
        x[k] = b[k] / a[k][k];
        for (i = k+1; i < n; i++) {
            for (j = k+1; j < n; j++) {
                a[i][j] -= a[i][k] * a[k][j];
            }
            b[i] = a[i][k] * x[k];
            a[i][k] = 0;
        }
    }
}
```

Complexity: $T = O(n^3)$

Naive Gaussian Elimination (cont.)



Gaussian Elimination with Partial-Pivoting

The naive GE algorithm runs into trouble when a pivot element is zero. (In fact, it runs into trouble whenever a pivot element is “small.”)

Partial-Pivoting — Search the column below the diagonal and find the largest element in absolute magnitude; perform row interchange to bring this element to the diagonal.

Example:

```
Initial:                                => Pivoting:
    2.0   1.0   3.0   4.0  29.0          5.0   2.0   0.0   1.0  13.0
    1.0   0.0   0.0   1.0   5.0          1.0   0.0   0.0   1.0   5.0
    3.0   1.0   1.0   0.0   8.0          3.0   1.0   1.0   0.0   8.0
    5.0   2.0   0.0   1.0  13.0          2.0   1.0   3.0   4.0  29.0

=> Elimination:                        =>
    5.0   2.0   0.0   1.0  13.0
    0.0 -0.4   0.0   0.8   2.4      Continue for other columns ...
    0.0 -0.2   1.0 -0.6   0.2
    0.0   0.2   3.0   3.6  23.8
```

Parallel GE with 1-D Partitioning

P_0	1	(0,1)	(0,2)	(0,3)	(0,4)	(0,5)	(0,6)	(0,7)
P_1	0	1	(1,2)	(1,3)	(1,4)	(1,5)	(1,6)	(1,7)
P_2	0	0	1	(2,3)	(2,4)	(2,5)	(2,6)	(2,7)
P_3	0	0	0	(3,3)	(3,4)	(3,5)	(3,6)	(3,7)
P_4	0	0	0	(4,3)	(4,4)	(4,5)	(4,6)	(4,7)
P_5	0	0	0	(5,3)	(5,4)	(5,5)	(5,6)	(5,7)
P_6	0	0	0	(6,3)	(6,4)	(6,5)	(6,6)	(6,7)
P_7	0	0	0	(7,3)	(7,4)	(7,5)	(7,6)	(7,7)

(a) Computation:

(i) $A[k,j] := A[k,j]/A[k,k]$ for $k < j <$

(ii) $A[k,k] := 1$

P_0	1	(0,1)	(0,2)	(0,3)	(0,4)	(0,5)	(0,6)	(0,7)
P_1	0	1	(1,2)	(1,3)	(1,4)	(1,5)	(1,6)	(1,7)
P_2	0	0	1	(2,3)	(2,4)	(2,5)	(2,6)	(2,7)
P_3	0	0	0	1	$\frac{1}{2}$	$\frac{1}{3}$	$\frac{1}{4}$	$\frac{1}{5}$
P_4	0	0	0	0	$\frac{1}{4}$	$\frac{1}{3}$	$\frac{1}{2}$	$\frac{1}{6}$
P_5	0	0	0	0	$\frac{1}{5}$	$\frac{1}{4}$	$\frac{1}{3}$	$\frac{1}{5}$
P_6	0	0	0	0	$\frac{1}{6}$	$\frac{1}{5}$	$\frac{1}{4}$	$\frac{1}{6}$
P_7	0	0	0	0	$\frac{1}{7}$	$\frac{1}{6}$	$\frac{1}{5}$	$\frac{1}{7}$

(b) Communication:

One-to-all broadcast of row $A[k, *]$

P_0	1	(0,1)	(0,2)	(0,3)	(0,4)	(0,5)	(0,6)	(0,7)	
P_1	0	1	(1,2)	(1,3)	(1,4)	(1,5)	(1,6)	(1,7)	
P_2	0	0	1	(2,3)	(2,4)	(2,5)	(2,6)	(2,7)	
P_3	0	0	0	1	(3,4)	(3,5)	(3,6)	(3,7)	
P_4	0	0	0		(4,3)	(4,4)	(4,5)	(4,6)	(4,7)
P_5	0	0	0		(5,3)	(5,4)	(5,5)	(5,6)	(5,7)
P_6	0	0	0		(6,3)	(6,4)	(6,5)	(6,6)	(6,7)
P_7	0	0	0		(7,3)	(7,4)	(7,5)	(7,6)	(7,7)

(c) Computation:

(i) $A[i,j] := A[i,j] - A[i,k] \times A[k,j]$
for $k < i < n$ and $k < j < n$

(ii) $A[i,k] := 0$ for $k < i < n$

Parallel GE with 1-D Partitioning

- ▶ Assign one row to each process.
- ▶ Execute the outer loop sequentially.
- ▶ At iteration $k+1$, process P_k either broadcast or shift its row to processes P_{k+1}, \dots, P_{n-1} .
- ▶ Each process performs local computation.

Complexity:

$$T_p = \frac{3}{2}n(n-1) + t_s n \log n + \frac{1}{2}t_w n(n-1) \log n$$

Improvements

Problems with the Simple Algorithm — The partition is fine-grained, and towards the end, the active region of the matrix is shrinking towards lower right corner; implies that processes fall idle.

Solutions: block and cyclic partitions.

P_0	1	(0,1)	(0,2)	(0,3)	(0,4)	(0,5)	(0,6)	(0,7)
	0	1	(1,2)	(1,3)	(1,4)	(1,5)	(1,6)	(1,7)
P_1	0	0	1	(2,3)	(2,4)	(2,5)	(2,6)	(2,7)
	0	0	0	1	(3,4)	(3,5)	(3,6)	(3,7)
P_2	0	0	0	(4,3)	(4,4)	(4,5)	(4,6)	(4,7)
	0	0	0	(5,3)	(5,4)	(5,5)	(5,6)	(5,7)
P_3	0	0	0	(6,3)	(6,4)	(6,5)	(6,6)	(6,7)
	0	0	0	(7,3)	(7,4)	(7,5)	(7,6)	(7,7)

P_0	1	(0,1)	(0,2)	(0,3)	(0,4)	(0,5)	(0,6)	(0,7)
	0	1	(1,2)	(1,3)	(1,4)	(1,5)	(1,6)	(1,7)
P_1	0	0	1	(2,3)	(2,4)	(2,5)	(2,6)	(2,7)
	0	0	0	(3,3)	(3,4)	(3,5)	(3,6)	(3,7)
P_2	0	0	0	(4,3)	(4,4)	(4,5)	(4,6)	(4,7)
	0	0	0	(5,3)	(5,4)	(5,5)	(5,6)	(5,7)
P_3	0	0	0	(6,3)	(6,4)	(6,5)	(6,6)	(6,7)
	0	0	0	(7,3)	(7,4)	(7,5)	(7,6)	(7,7)

(a) Block 1-D mapping

P_0	1	(0,1)	(0,2)	(0,3)	(0,4)	(0,5)	(0,6)	(0,7)
	0	0	0	(4,3)	(4,4)	(4,5)	(4,6)	(4,7)
P_1	0	1	(1,2)	(1,3)	(1,4)	(1,5)	(1,6)	(1,7)
	0	0	0	(5,3)	(5,4)	(5,5)	(5,6)	(5,7)
P_2	0	0	1	(2,3)	(2,4)	(2,5)	(2,6)	(2,7)
	0	0	0	(6,3)	(6,4)	(6,5)	(6,6)	(6,7)
P_3	0	0	0	(3,3)	(3,4)	(3,5)	(3,6)	(3,7)
	0	0	0	(7,3)	(7,4)	(7,5)	(7,6)	(7,7)

(b) Cyclic 1-D mapping

Parallel GE with Partial-Pivoting

- ▶ Distribute the matrix as complete rows; each process stores approximately n/p rows of the matrix.
- ▶ Processes collectively decide on which two rows need to be swapped to do the partial pivoting.
- ▶ Two processes do a send/recv to each other to do the swap. (Alternatively, we can keep track of which row is which through an indirection array.)
- ▶ The pivot row is broadcast to all processes.
- ▶ In each process, we loop over those rows below the pivot row and apply the elimination step (an `saxby()` operation).
- ▶ Repeat until we hit bottom.

Parallel GE with 2-D Partitioning

Each processor gets a 2D block of the matrix.

Steps:

- ▶ Broadcast the “active” column along the rows.
- ▶ Prepare the pivot row concurrently.
- ▶ Broadcast the pivot row along the columns.
- ▶ Perform the elimination step concurrently.

1	(0,1)	(0,2)	(0,3)	(0,4)	(0,5)	(0,6)	(0,7)
0	1	(1,2)	(1,3)	(1,4)	(1,5)	(1,6)	(1,7)
0	0	1	(2,3)	(2,4)	(2,5)	(2,6)	(2,7)
0	0	0	$\begin{matrix} \vdots \\ (3,3) \end{matrix}$	(3,4)	(3,5)	(3,6)	(3,7)
0	0	0	$\begin{matrix} \vdots \\ (4,3) \end{matrix}$	(4,4)	(4,5)	(4,6)	(4,7)
0	0	0	$\begin{matrix} \vdots \\ (5,3) \end{matrix}$	(5,4)	(5,5)	(5,6)	(5,7)
0	0	0	$\begin{matrix} \vdots \\ (6,3) \end{matrix}$	(6,4)	(6,5)	(6,6)	(6,7)
0	0	0	$\begin{matrix} \vdots \\ (7,3) \end{matrix}$	(7,4)	(7,5)	(7,6)	(7,7)

(a) Rowwise broadcast of $A[i,k]$
for $(k-1) < i < n$

1	(0,1)	(0,2)	(0,3)	(0,4)	(0,5)	(0,6)	(0,7)
0	1	(1,2)	(1,3)	(1,4)	(1,5)	(1,6)	(1,7)
0	0	1	(2,3)	(2,4)	(2,5)	(2,6)	(2,7)
0	0	0	(3,3)	(3,4)	(3,5)	(3,6)	(3,7)
0	0	0	(4,3)	(4,4)	(4,5)	(4,6)	(4,7)
0	0	0	(5,3)	(5,4)	(5,5)	(5,6)	(5,7)
0	0	0	(6,3)	(6,4)	(6,5)	(6,6)	(6,7)
0	0	0	(7,3)	(7,4)	(7,5)	(7,6)	(7,7)

(b) $A[k,j] := A[k,j]/A[k,k]$
for $k < j < n$

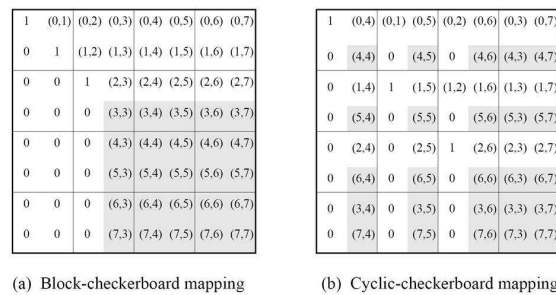
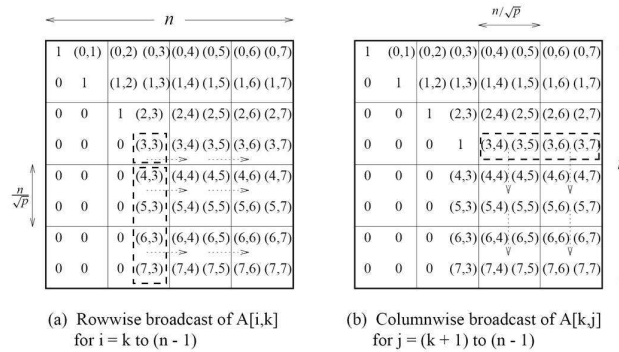
1	(0,1)	(0,2)	(0,3)	(0,4)	(0,5)	(0,6)	(0,7)
0	1	(1,2)	(1,3)	(1,4)	(1,5)	(1,6)	(1,7)
0	0	1	(2,3)	(2,4)	(2,5)	(2,6)	(2,7)
0	0	0	1	$\begin{matrix} \vdots \\ (3,4) \end{matrix}$	(3,5)	(3,6)	(3,7)
0	0	0	$\begin{matrix} \vdots \\ (4,4) \end{matrix}$	(4,4)	(4,5)	(4,6)	(4,7)
0	0	0	$\begin{matrix} \vdots \\ (5,4) \end{matrix}$	(5,4)	(5,5)	(5,6)	(5,7)
0	0	0	$\begin{matrix} \vdots \\ (6,4) \end{matrix}$	(6,4)	(6,5)	(6,6)	(6,7)
0	0	0	$\begin{matrix} \vdots \\ (7,4) \end{matrix}$	(7,4)	(7,5)	(7,6)	(7,7)

(c) Columnwise broadcast of $A[k,j]$
for $k < j < n$

1	(0,1)	(0,2)	(0,3)	(0,4)	(0,5)	(0,6)	(0,7)
0	1	(1,2)	(1,3)	(1,4)	(1,5)	(1,6)	(1,7)
0	0	1	(2,3)	(2,4)	(2,5)	(2,6)	(2,7)
0	0	0	1	(3,4)	(3,5)	(3,6)	(3,7)
0	0	0	(4,3)	(4,4)	(4,5)	(4,6)	(4,7)
0	0	0	(5,3)	(5,4)	(5,5)	(5,6)	(5,7)
0	0	0	(6,3)	(6,4)	(6,5)	(6,6)	(6,7)
0	0	0	(7,3)	(7,4)	(7,5)	(7,6)	(7,7)

(d) $A[i,j] := A[i,j] - A[i,k] \times A[k,j]$
for $k < i < n$ and $k < j < n$

Block and Cyclic 2-D Partitions



Solving Large, Sparse Linear Systems

Large sparse matrices often appear in scientific or engineering applications when solving partial differential equations. As an example, consider the Laplace Equation:

$$\phi_{i,j} = \frac{1}{4}(\phi_{i+1,j} + \phi_{i-1,j} + \phi_{i,j+1} + \phi_{i,j-1}), \quad 0 < i,j < n$$

It describes a computation over an $n \times n$ mesh. $\phi_{i,j}$ denotes the value at the mesh point $[i,j]$.

```
# # # ... # # # ..#
# . . . . . . . . . #
. . . . . . . . . . .
# . . . . . n . . . #
# . . . . w x e . . #
# . . . . . s . . . #
. . . . . . . . . . .
. . . . . . . . . . .
# . . . . . . . . . #
# # # ... # # # ..#
```

- The value at point $[i,j]$ is related to the values of its four neighbors.
- The values at the four boarder lines are typically fixed.

Solving the Laplace Equation

The Laplace Equation is a linear system with n^2 unknowns. It can be turned into the standard $\mathbf{Ax} = \mathbf{b}$ form:

Let $l = i \cdot n + j$, then the system becomes:

$$\phi_{l+n} + \phi_{l-n} + \phi_{l+1} + \phi_{l-1} - 4\phi_l = 0, \quad 0 < l < n^2$$

which can also be expressed in matrix form:

$$\begin{array}{cccccccc} -4 & 1 & 0 & 0 & \dots & 1 & 0 & \dots \\ 1 & -4 & 1 & 0 & \dots & 0 & 1 & \dots \\ 0 & 1 & -4 & 1 & \dots & 0 & 0 & \dots \\ 0 & 0 & 1 & -4 & \dots & \dots & \dots & \dots \\ \dots & \dots & \dots & \dots & \dots & \dots & \dots & \dots \\ 1 & 0 & 0 & \dots & \dots & \dots & \dots & \dots \\ 0 & 1 & 0 & \dots & \dots & \dots & \dots & \dots \\ \dots & \dots & \dots & \dots & \dots & \dots & \dots & \dots \\ \dots & \dots & \dots & \dots & \dots & \dots & \dots & \dots \\ \dots & \dots & \dots & \dots & \dots & \dots & \dots & \dots \end{array}$$

- The matrix is very large, yet sparse.
- It can be solved by using Gaussian Elimination, but the cost would be very high: $O(n^6)$.

A better approach is to directly solve the system over the original $n \times n$ mesh domain.

Jacobi Relaxation Algorithm

It is an iterative algorithm:

- ▶ Cycle through the mesh points, compute a new value for each point using the average of the four-neighboring “old” values.

$$\phi_{i,j}^{(t+1)} = \frac{1}{4}(\phi_{i+1,j}^{(t)} + \phi_{i-1,j}^{(t)} + \phi_{i,j+1}^{(t)} + \phi_{i,j-1}^{(t)}), \quad 0 < i, j < N$$

- ▶ Once all the new values are found, replace old values with new ones.
- ▶ Repeat until the difference between old and new is small enough.

Parallelization:

- ▶ For the distributed-memory case, simply partition the mesh in both dimensions; each processor will handle a $\frac{n}{\sqrt{p}} \times \frac{n}{\sqrt{p}}$ sub-mesh. Communications are only required on the “peripheral” of the mesh region in each processor.

Jacobi Relaxation Algorithm (cont.)

```
int jacobi(int n, double x[n][n], double epsilon) {
    double xnew[n][n], delta;
    int i, j, cnt = 0;
    do {
        cnt++;
        delta = 0.0;
        for (i = 1; i < n-1; i++) {
            for (j = 1; j < n-1; j++) {
                xnew[i][j] = (x[i-1][j] + x[i][j-1]
                             + x[i+1][j] + x[i][j+1]) / 4.0;
                delta = fmax(delta, fabs(xnew[i][j] - x[i][j]));
            }
        }
        for (i = 1; i < n-1; i++) {
            for (j = 1; j < n-1; j++) {
                x[i][j] = xnew[i][j];
            }
        }
    } while (delta > epsilon);
    return cnt;
}
```

Gauss-Seidel Algorithm

While simple, the Jacobi relaxation algorithm has two drawbacks, (1) It needs buffers to store old mesh point values; and (2) Its convergence speed can be slow.

Gauss-Seidel algorithm is an improvement over the Jacobi algorithm. It's like Jacobi, except that mesh points are updated "in-place," overwriting the old value.

The order of mesh point updates does not really matter. Here are two possible orders:

$$\begin{aligned}\phi_{i,j}^{(t+1)} &= \frac{1}{4}(\phi_{i+1,j}^{(t)} + \phi_{i-1,j}^{(t+1)} + \phi_{i,j+1}^{(t)} + \phi_{i,j-1}^{(t+1)}), & 0 < i,j < N \\ \phi_{i,j}^{(t+1)} &= \frac{1}{4}(\phi_{i+1,j}^{(t+1)} + \phi_{i-1,j}^{(t)} + \phi_{i,j+1}^{(t+1)} + \phi_{i,j-1}^{(t)}), & 0 < i,j < N\end{aligned}$$

- Gauss-Seidel algorithm does not need buffers and has a better convergence speed (since newer values are used in all updates).

Parallelizing Gauss-Seidel

Due to the “in-place” updates, race condition may happen among multiple reads and a single write to the same memory location. But this is not really a problem for an iterative algorithm — using an old value vs. using a new value affects only the convergency rate.

The following are the common parallel approaches:

- ▶ *Natural index ordering* (allow race condition)

- ▶ *Red-black ordering* (avoid race condition)

- Denote alternating points as “red” and “black”.
- Repeat following two steps until convergence:
 - update all red points (in any order)
 - update all black points (in any order)

```

r  b  r  b  r  b
b  r  b  r  b  r
r  b  r  b  r  b
b  r  b  r  b  r
r  b  r  b  r  b
b  r  b  r  b  r

```

- ▶ *Wavefront ordering* (avoid race condition)

- Proceed along diagonal line.
- Repeat until success.

```

1  2  3  4  ..  ..
2  3  4  5  ..  ..
3  4  5  ..  ..  ..
4  5  ..  ..  ..  ..
5  ..  ..  ..  ..  ..
.  .  .  .  .  .

```

Multi-Grid Methods

Observations — Iterative algorithms converge to solutions faster on coarser grids than on finer grids. Iterative algorithms converge quicker if the initial approx. of the values of the variables are good.

Algorithm:

- ▶ Begin with the original problem, where the solution is defined (and desired) on an $n \times n$ mesh.
- ▶ (*Projection*) Coarsen the problem by several powers of 2: $n \rightarrow n/2^m$. Boundary values need to be averaged down to the coarser mesh.
- ▶ (*Relaxation*) Solve the coarse version problem using any iterative solver. It should go fast since the mesh is small.
- ▶ (*Interpolation*) Boost up the problem by a factor of 2, interpolating field points. What was one mesh point turns into 4 mesh points.
- ▶ Re-run relaxation on this problem.
- ▶ Repeat *Relaxation/Interpolation* steps until back to original problem.