

Data Parallel Programming

Jingke Li

Portland State University

Programming Vector/SIMD Computers

Vector computers and SIMD computers have a common programming interface:

- ▶ a single stream of instructions
- ▶ each instruction specifies *uniform* operations over a set of data
- ▶ the operations are carried out in a synchronized step.

This style of computation is called *data parallel*.

Two Programming Approaches:

- ▶ *Regular programs + smart compiler*
 - compiler detects data parallelism in user program; converts them to vector form
- ▶ *Explicit data-parallel programs*
 - requires language support

The Compiler Approach

Source code:

```
for i = 1 to n do
  a[i] = b[i] + c[i]
  if a[i]>0 then
    b[i] = b[i] + 1
  endif
endfor
```

Vector code:

```
for i = 1 to n by 64 do
  VL = max(n-i+1,64)
  vfetch b[i] -> v1
  vfetch c[i] -> v2
  vadd v1+v2 -> v3
  vstore v3 -> a[i]
  loadi #0 -> r1
  vcmp v3, r1 -> m1
  loadi #1 -> r2
  vaddc v1+r1 -> v1, m1
  vstorec v1 -> b[i], m1
endfor
```

Key compiler techniques:

- ▶ data dependence analysis
- ▶ loop transformations

Explicit Programming Approach

Use simple, high-level constructs to specify data parallelism — uniform computations performed on all elements of a data domain.

- ▶ *Key constructs*: array operations and parallel loops
- ▶ Synchronizations are implicitly
- ▶ Communications are not allowed

Languages:

- ▶ Fortran 90/95
- ▶ OpenCL and Cuda
- ▶ Chapel

Fortran 90/95

Fortran 90 is a major extension to Fortran 77 for

- ▶ data-parallel style parallel programming
- ▶ language modernization (e.g. dynamic memory management, recursive functions, derived types, pointers, bit functions, etc.)

Fortran 77:

```
real sums
dimension sums(2,3)
do 20 j = 1, 3
  do 10 i = 1, 2
    sums(i,j) = 0.0
  10 continue
20 continue
```

Fortran 90:

```
real,dimension(2,3):: sums = 0.0
```

Whole Array Operations

```
program simple
integer a, b, c, n, max
parameter (n=5)
dimension a(n), b(n), c(n)
data a /1,2,3,4,5/ ! array initialization

b = 2 ! scalar expansion
c = a**2 + b**2 ! whole array operations
print *, 'array c contains:'
print *, c
max = maxval(c)
print *, 'the largest value in c is ', max

stop
end
```

Array Section Operations

Triplet Notation: (lbound : ubound : stride)

```
integer,dimension(5,5):: array
integer,dimension(5):: row1,row2,row3,col5
integer,dimension(3,3):: inner
integer,dimension(25):: linear
integer,dimension(5):: diagonal

Row1 = array(1,1:5:1)
Row2 = array(2,1:5)
Row3 = array(3,:)
Col5 = array(:,5)
inner = array(2:4,2:4)
linear = pack(array, .true.)
diagonal = linear(1:25:6)
```

Array Functions

- Array attributes: lbound, ubound, shape, size

```
real,dimension(-1:2,-3:4):: array
print *, lbound(array) ! -1 -3
print *, ubound(array) ! 2 4
print *, shape(array) ! 4 8
print *, size(array) ! 32
```

- Rearranging elements: reshape, pack, unpack

```
integer,dimension(3,3):: array
= reshape( (/ 1,4,7,2,5,8,3,6,9 /), (/ 3,3 /) )
integer,dimension(9):: vector = (/ (i,i=1,9) /)
logical,dimension(3,3):: filter = .true.
integer:: missing = 0
filter(:,2) = .false. ! [1 4 7] [1 0 7]
filter(2,:) = .false. ! [2 5 8] -> [0 0 0]
array = unpack(vector, filter, missing) ! [3 6 9] [3 0 9]
```

- Intrinsic Functions: (Built-in parallel functions)
 - sum, product, max, min, any, all, count, matmul, spread, etc.

Conditional Operations

- The original “if” statement:
 - cond is a scalar boolean expression

```
integer,dimension(10,10):: a, b, c
if (k > 5) then
    c = b/a
end if
```

- The new parallel “where” statement:
 - cond is an array boolean expression

```
integer,dimension(10,10):: a, b, c
where (a .ne. 0)
    c = b/a
elsewhere
    c = -1
end where
```

Prime-Finder Program Comparison

Fortran 77:

```
program prime
integer i, j, n
parameter (n=999)
logical primes(n)

do i = 1, n
    primes(i) = .true.
end do
primes(1) = .false.
do i = 1, int(sqrt(real(n)))
    if (primes(i)) then
        do j = i+1, n, i
            primes(j) = .false.
        end do
    end if
end do
do i = 1, n
    if (primes(i)) print *, i
end do
end
```

Fortran 90:

```
program rimes
integer i, n, next
parameter (n=999)
logical,dimension(n):: primes
integer,dimension(n):: ident = (/ (i,i=1,n) /)

primes = .true.
primes(1) = .false.

next = 2
do while (next .lt. int(sqrt(real(n))))
    primes(next*2:n:next) = .false.
    next = minval(ident(next+1:n),
                    primes(next+1:n))
end do

print *, "Number of primes:", count(primes)
do i = 1, n
    if (primes(i)) print *, i
end do
end program
```

Fortran 90 Prime-Finder Example

- Initialization (assume n=10):

```
n = 10
primes = (f,t,t,t,t,t,t,t,t,t)
ident = (1,2,3,4,5,6,7,8,9,10)
next = 2
```

- 1st iteration:

```
primes(4:10:2) = .false.
primes = (f,t,t,f,t,f,t,f,t,f)
next = minval((3,4,5,6,7,8,9,10), (t,f,t,f,t,f,t,f))
⇒ next = 3
```

- 2nd iteration:

```
primes(6:10:3) = .false.
primes = (f,t,t,f,t,f,t,f,f,f)
next = minval((4,5,6,7,8,9,10), (f,t,f,t,f,f,f,f))
⇒ next = 5 (> loop bound)
```

The Fortran 95 Forall Statement

The intent of the Forall statement is that its “iterations” can be executed *concurrently*. For example,

```
forall (i = 1:n)
    a(i) = b(i) + c(i)
end forall
```

The n iterations imply n concurrent threads. At the end of the loop, there is an implicit *barrier* synchronization to synchronize all n threads.

However, defining its semantics is not as easy as it appears. Consider:

```
forall (i = 2:n)
    a(i) = b(i) + a(i-1)
end forall
```

```
forall (i = 1:n-1)
    a(i) = b(i) + c(i)
    a(i+1) = a(i) + a(i+1)
end forall
```

Both contain inter-iteration dependencies. The second example is further complicated by having multiple statements in the body of the loop. How should they be executed?

Forall Statement Semantics

- ▶ Only assignment statements, where statements, and nested forall statements are allowed inside a forall statement.
- ▶ All reads from a given assignment statement, in all iterations, must occur before any write to the left-hand side, in any iteration.
- ▶ The writes of the left-hand side must occur before any reads in the following assignment statement.

Back to the Examples

```
forall (i = 2:n)
  a(i) = b(i) + a(i-1)
end forall
```

- ▶ The value of $a(i-1)$ used in iteration i is the *old* value from before the loop execution starts; it is not the value produced in iteration $i-1$.

```
forall (i = 1:n-1)
  a(i) = b(i) + c(i)
  a(i+1) = a(i) + a(i+1)
end forall
```

- ▶ The value of $a(i)$ used in the second statement is the *new* value produced by the first statement of the same iteration.
- ▶ The value of $a(i+1)$ used in the second statement is the *old* value from before the loop execution starts.

Sequential Execution of Forall Statement

The forall statements are best implemented on a vector or a SIMD machine. However, there will be times they need to be executed sequentially, for instance, when there is not enough processors to match the number of iterations, or when we want to test and debug a parallel program on a sequential machine.

When running on a single processor, a forall statement can be expensive, because in general, it cannot be simply turned into a do loop. In other words, the following two loops may produce different results:

```
forall (i = 2:n)
  a(i) = b(i) + a(i-1)
end forall
```

```
do i = 2 to n
  a(i) = b(i) + a(i-1)
end do
```

The correct implementation calls for buffering a copy of the old values of a .