

Pthreads

Jingke Li

Portland State University

Pthreads

The standard thread package on Unix/Linux systems. To use it on Linux,

- ▶ Include the header file `<pthread.h>` in your program.
- ▶ Compile your program with the “`-pthread`” flag.

Useful Routines:

```
pthread_create(&thread, attributes, func, arg);
pthread_exit(status);
pthread_join(thread, &status); // sync back a thread
pthread_cancel(thread);       // force termination a thread
pthread_self();               // return the self thread
pthread_equal(thread1, thread2); // compare two threads
sched_yield();                // yield to another thread
```

- The `pthread_join()` routine is for synchronizing a *terminated* thread back to its parent.
- Not all threads are joinable — a thread can be created as *detached*; when it terminates, it is destroyed and its resource released.

Synchronizations in Pthreads

► Mutex Locks:

```
pthread_mutex_t mutex;  
  
pthread_mutex_init(&mutex, attributes);  
pthread_mutex_lock(&mutex);  
pthread_mutex_trylock(&mutex);  
pthread_mutex_destroy();
```

► Condition Variables:

```
pthread_cond_t cond;  
  
pthread_cond_init(&cond, attributes);  
pthread_cond_signal(cond);  
pthread_cond_broadcast(cond);  
pthread_cond_wait(cond, mutex);
```

Attributes

Various attributes can be set for threads, locks, and condition variables. For normal cases, the default values work just fine.

- For threads:
 - scheduling policy, stack size, detached state
- For mutexes:
 - *normal* — only a single thread is allowed to lock it; if a threads tries to lock it twice a deadlock occurs.
 - *recursive* — a thread can lock the mutex multiple time; each successive lock/unlock increments/decrements a counter; another thread can lock a mutex only if its counter is zero.
- For condition variables:
 - can enable cross-process sharing

A Simple Example

```
#include <pthread.h>

void child(long id) {
    printf("Child %ld\n", id);
}

main() {
    long i=1, j=2;
    pthread_t thread1, thread2;
    pthread_create(&thread1, NULL,
                  (void*) child, (void*)i);
    pthread_create(&thread2, NULL,
                  (void*) child, (void*)j);
    printf("Parent: waiting for child 1\n");
    pthread_join(thread1, NULL);
    printf("Parent: waiting for child 2\n");
    pthread_join(thread2, NULL);
    printf("Parent: done\n");
}
```

Output:

```
Parent: waiting for child 1
Child 1
Child 2
Parent: waiting for child 2
Parent: done
```

Mutexes

Locks are implemented in Pthreads with “mutex” variables.

- ▶ To use a mutex, first it must be declared and initialized:

```
pthread_mutex_t mutex;
pthread_mutex_init(&mutex, NULL);
```

NULL specifies a default attribute for the mutex.

- ▶ A critical section can then be protected with the mutex:

```
pthread_mutex_lock(&mutex);
<critical section>
pthread_mutex_unlock(&mutex);
```

- If a thread reaches a mutex lock and finds it locked, it will wait for the lock to open.
- If more than one thread is waiting for the lock to open when it opens, the system will select one thread to be allowed to proceed.

Mutexes (cont.)

- ▶ Only the thread that locks a mutex can unlock it.
- ▶ Pthreads offers a routine that can test whether a lock is locked without blocking the thread:

```
pthread_mutex_trylock()
```

- It will lock an unlocked mutex and return 0 or will return with EBUSY if the mutex is locked — useful in overcoming deadlock.

- ▶ A mutex can be destroyed with

```
pthread_mutex_destroy()
```

Condition Variables

- ▶ Condition variables also need to be initialized before use:

```
pthread_cond_t cond;  
pthread_cond_init(&cond, NULL);
```

- ▶ Each condition variable must be associated with a mutex, since the checking and setting of the condition must be done inside a critical section. The “wait” routine, in particular, takes a mutex as one of its arguments:

```
pthread_cond_wait(cond, mutex);
```

- ▶ Signals that are sent out by “signal” or “broadcast” routines are not remembered, which means that threads must already be waiting for a signal to receive it.

Condition Variable Example

Decrement a count; if value reaches 0, send a signal.

```
counter() {
    pthread_mutex_lock(&mutex);
    c--;
    if (c == 0)
        pthread_cond_signal(cond);
    pthread_mutex_unlock(&mutex);
}
```

Compare with two other versions:

```
counter_v2() {
    pthread_mutex_lock(&mutex);
    c--;
    if (c == 0)
        pthread_cond_broadcast(cond);
    pthread_mutex_unlock(&mutex);
}
```

```
counter_v3() {
    pthread_mutex_lock(&mutex);
    c--;
    pthread_mutex_unlock(&mutex);
    if (c == 0)
        pthread_cond_signal(cond);
}
```

Condition Variable Example (cont.)

```
action() {
    pthread_mutex_lock(&mutex);
    while (c <> 0)
        pthread_cond_wait(cond, mutex);
    pthread_mutex_unlock(&mutex);
    take_action();
}
```

Compare with the following version:

```
action_v2() {
    pthread_mutex_lock(&mutex);
    if (c <> 0)
        pthread_cond_wait(cond, mutex);
    pthread_mutex_unlock(&mutex);
    take_action();
}
```

Array-Sum Example

```
#include <pthread.h>

int arraySize = 1000;           // default array size
int numThreads = 10;           // default number of threads
int *array;                     // shared array
int sum = 0, idx = 0;           // global sum and idx
pthread_mutex_t sumLock;

int main(int argc, char **argv) {
    pthread_t thread[numThreads];
    array = init_array(arraySize); // initialize array
    pthread_mutex_init(&sumLock, NULL); // initialize mutex

    for (long k = 0; k < numThreads; k++) { // create threads
        pthread_create(&thread[k], NULL, (void*)slave, (void*)k);
    }
    for (long k = 0; k < numThreads; k++) { // join threads
        pthread_join(thread[k], NULL);
    }
    printf("The sum of 1 to %i is %d\n", arraySize, sum);
}
```

Array-Sum Example (cont.)

```
int *init_array(int size) {
    int *array = (int *) malloc(sizeof(int) * size);
    for (int i = 0; i < size; i++)
        array[i] = i + 1;
    return array;
}

void slave(long tid) {
    printf("Thread %ld started\n", tid);

    int i, psum = 0;
    do {
        pthread_mutex_lock(&sumLock); // read and increment idx
        i = idx++;
        pthread_mutex_unlock(&sumLock);
        if (i < arraySize) // add one array element
            psum += array[i];
    } while (i < arraySize);
    pthread_mutex_lock(&sumLock); // add local psum to global sum
    sum += psum;
    pthread_mutex_unlock(&sumLock);
}
```

Array-Sum Version 2

Show where threads are executed.

```
#define _GNU_SOURCE
#include <pthread.h>
#include <sched.h>           // for getting cpu id

void slave(long tid) {
    printf("Thread %ld started on %d\n", tid, sched_getcpu());

    ...
}
```

Question: Can we have more control over where threads execute?

Array-Sum Version 3

Control where threads are executed.

```
#define _GNU_SOURCE
#include <pthread.h>
#include <sched.h>           // for getting cpu id
#include <unistd.h>          // for getting nprocs

int main(int argc, char **argv) {
    pthread_t thread[numThreads];
    array = init_array(arraySize);           // initialize array
    pthread_mutex_init(&sumLock, NULL);       // initialize mutex
    int nprocs = sysconf(_SC_NPROCESSORS_ONLN);
    cpu_set_t cpuset;
    int cid = 0;
    for (long k = 0; k < numThreads; k++) {    // create threads
        pthread_create(&thread[k], NULL, (void*)slave, (void*)k);
        CPU_ZERO(&cpuset);
        CPU_SET(cid++ % nprocs, &cpuset);
        pthread_setaffinity_np(thread[k], sizeof(cpu_set_t), &cpuset);
    }
    ...
}
```

Array-Sum Version 4

Add command-line arguments for parameter configurations.

```
int arraySize;           // array size, given by user
int numThreads = 1;      // default number of threads

int main(int argc, char **argv) {
    if (argc < 2) {
        printf ("Usage: ./arraysum4 <arraySize> [<numThreads>]\n");
        exit(0);
    } else if (argc > 2) {
        if ((numThreads=atoi(argv[2])) < 1) {
            printf ("<numThreads> must be greater than 0\n");
            exit(0);
        }
    }
    if ((arraySize=atoi(argv[1])) < 1) {
        printf ("<arraySize> must be greater than 0\n");
        exit(0);
    }
    ...
}
```

Producer-Consumer with Bounded Buffer

Problem Description:

- ▶ One producer, multiple consumers, and a bounded task queue.
- ▶ The producer creates tasks and adds them one by one to the end of the task queue. If the queue is full, it waits for new space to open up.
- ▶ Each consumer removes tasks one by one from the head of the task queue. If the queue is empty, it waits for new task to appear.

Programming Issues:

- ▶ Task and queue representations
- ▶ Threads creation and management
- ▶ Synchronization
- ▶ Termination

Producer-Consumer: Task and Queue Representations

```
typedef struct task_ task_t;
struct task_ {
    int val;           // each task holds an integer value
    task_t *next;      // and a pointer to next task
};

typedef struct queue_ queue_t;
struct queue_ {
    task_t *head;
    task_t *tail;
    int limit;         // queue size limit
    int length;        // current number of tasks
};
```

Supporting Routines:

```
task_t *create_task(int val) { ... }
queue_t *init_queue(int limit) { ... }
void add_task(queue_t *queue, task_t *task) { ... }
task_t *remove_task(queue_t *queue) { ... }
```

Producer-Consumer: Threads Creation and Management

```
int main(int argc, char **argv) {

    // create consumer threads
    pthread_t *threads =
        (pthread_t *) malloc(sizeof(pthread_t) * numConsumers);
    for (long k = 0; k < numConsumers; k++)
        pthread_create(&threads[k], NULL, (void*)consumer, (void*)k);

    // execute the producer code
    producer();

    // wait for consumer threads to terminate
    for (long k = 0; k < numConsumers; k++)
        pthread_join(threads[k], NULL);
}
```

Producer-Consumer: Synchronization

```
void producer() {
    while (<there is still task>) {
        <try to add a task to the queue>
        <wait if no space available>
    }
}

void consumer(long tid) {
    while (1) {
        <try to get a task from the queue>
        <wait if none available>
    }
}
```

Questions:

1. Other than the two waits, is there a need for additional synchronizations?
2. Who is/are responsible for sending signals to the waiting threads?

Producer-Consumer: Termination

```
void producer() {
    while (<there is still task>) {
        <try to add a task to the queue>
        <wait if no space available>
    }
}

void consumer(long tid) {
    while (1) {
        <try to get a task from the queue>
        <wait if none available>
    }
}
```

Questions:

1. Can the producer thread terminate on its own?
2. Can the consumer threads terminate on their own?
3. If not, what additional mechanism is needed?

Quicksort Program Framework

```
// A global array of size N contains the integers to be sorted.
// A global task queue is initialized with the sort range [0,N-1].
//
int main(int argc, char **argv) {
    // 1. read in command-line arguments, N and numThreads;
    // 2. initialize array, queue, and other shared variables

    // 3. create numThreads-1 worker threads, each executes a copy
    //     of the worker() routine
    for (long k = 0; k < numThreads-1; k++)
        pthread_create(&thread[k], NULL, (void*)worker, (void*)k);

    // 4. the main thread also runs a copy of the worker() routine;
    //     its copy has the last id, numThreads-1
    worker(numThreads-1);

    // 5. the main thread waits for worker threads to join back
    for (long k = 0; k < numThreads-1; k++)
        pthread_join(thread[k], NULL);

    // 6. verify the result
}
```

Quicksort Program Framework (cont.)

```
void worker(long wid) {
    while (<termination condition is not met>) {
        task = remove_task();
        quicksort(array, task->low, task->high);
    }
}

void quicksort(int *array, int low, int high, long wid) {
    // 1. find a pivot and partition the array into two segments
    int middle = partition(array, low, high);
    // 2. add the first segment to the task queue
    if (low < middle)
        <add task [low, middle-1] to queue>
    // 3. recursively sort on the second segment
    if (middle < high)
        quicksort(array, middle+1, high, wid);
}
```

Questions:

1. What synchronizations are needed?
2. What should the termination condition be?